

# Developing Microservices with Aggregates

Chris Richardson

Founder of Eventuate.io

Founder of the original CloudFoundry.com

Author of POJOs in Action

@crichtardson

chris@chrisrichardson.net

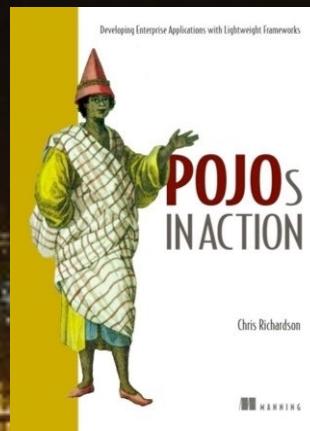
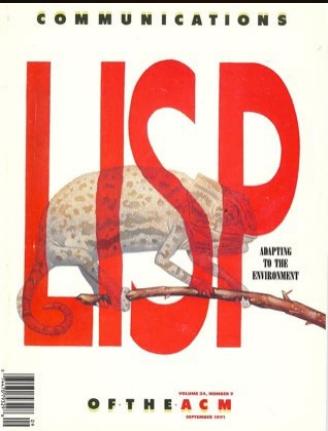
<http://eventuate.io>

SpringOne Platform

# Presentation goal

Show how  
Domain Driven Design Aggregates  
and Microservices  
are a perfect match

# About Chris



# About Chris

Consultant and trainer  
focusing on modern  
application architectures  
including microservices  
(<http://www.chrisrichardson.net/>)

# About Chris

Founder of a startup that is creating  
a platform that makes it easier for  
developers to write transactional  
microservices

(<http://eventuate.io>)



For more information

<http://learnmicroservices.io>

# Agenda

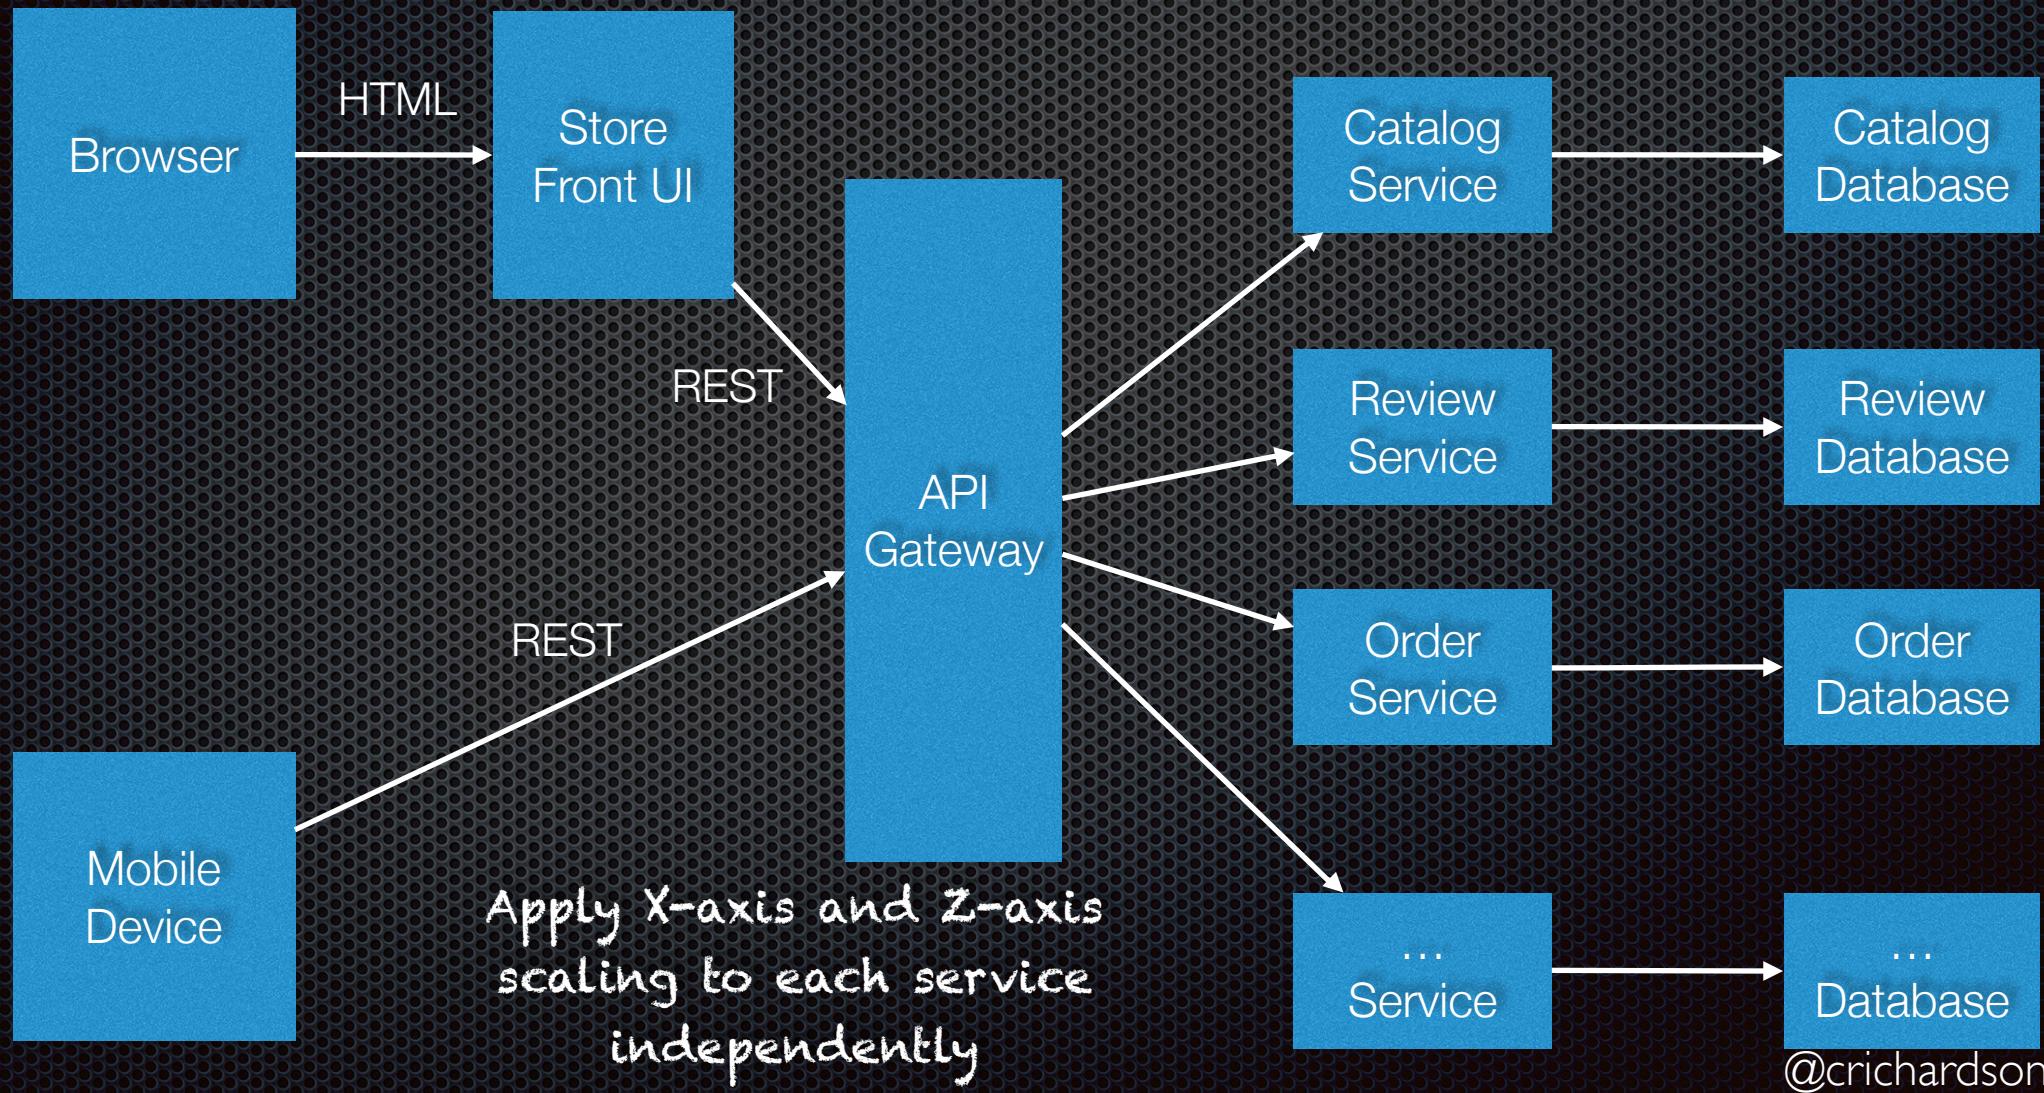
- The problem with Domain Models and microservices
- Overview of aggregates
- Maintaining consistency between aggregates
- Using event sourcing with Aggregates
- Example application

The Microservice architecture  
tackles complexity through  
modularization

Microservice  
=

Business capability

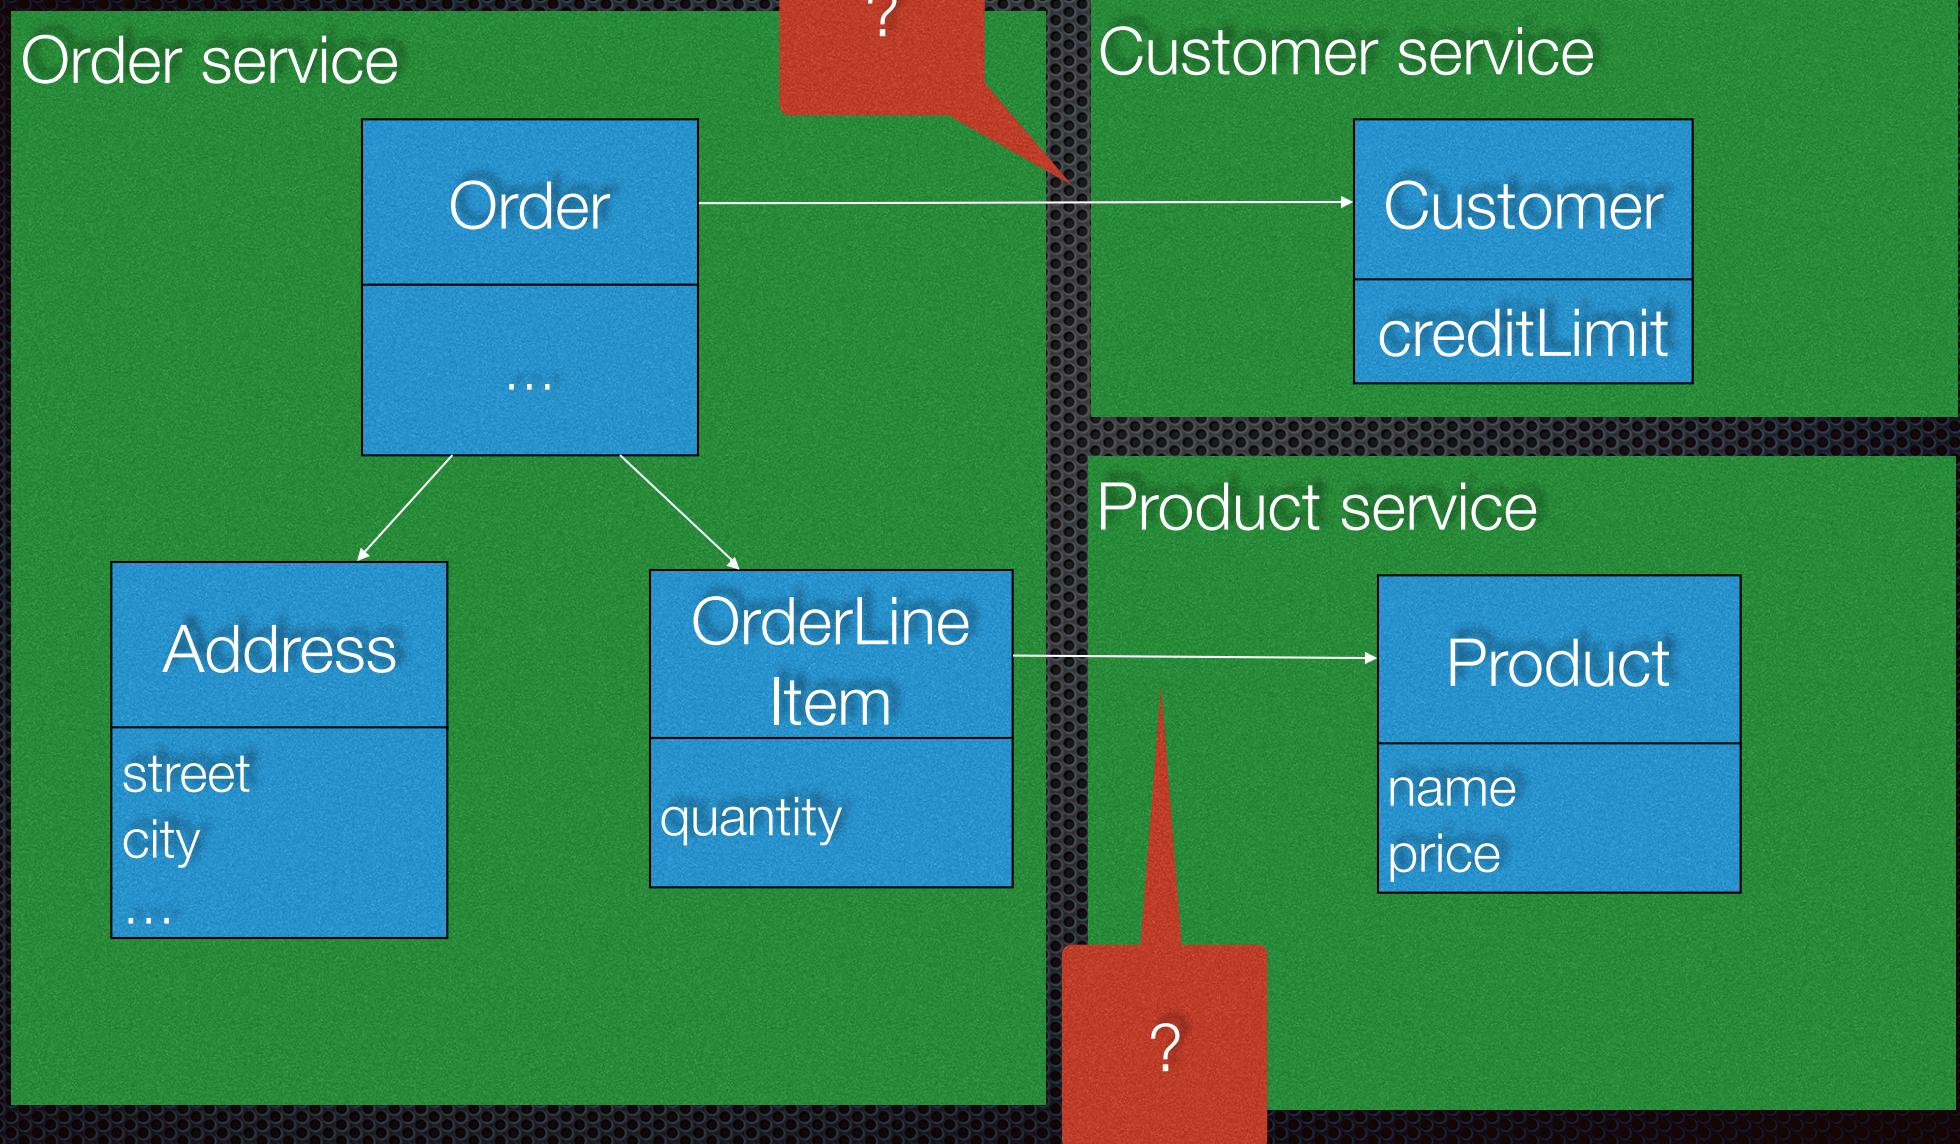
# Microservice architecture



Service boundaries  
enforce modularity

But there are challenges...

# Domain model = tangled web of classes



# Reliance on ACID transactions to enforce invariants

```
BEGIN TRANSACTION
```

```
...
```

```
SELECT ORDER_TOTAL  
FROM ORDERS WHERE CUSTOMER_ID = ?
```

```
...
```

```
SELECT CREDIT_LIMIT  
FROM CUSTOMERS WHERE CUSTOMER_ID = ?
```

```
...
```

```
INSERT INTO ORDERS ...
```

```
...
```

```
COMMIT TRANSACTION
```

Simple and  
ACID 😊

# But it violates encapsulation...

```
BEGIN TRANSACTION
```

```
...
```

```
SELECT ORDER_TOTAL
```

```
FROM ORDERS WHERE CUSTOMER_ID = ?
```

```
...
```

```
SELECT CREDIT_LIMIT
```

```
FROM CUSTOMERS WHERE CUSTOMER_ID = ?
```

```
...
```

```
INSERT INTO ORDERS ...
```

```
...
```

```
COMMIT TRANSACTION
```

Private to the  
Order Service

Private to the  
Customer Service

# .. and requires 2PC

```
BEGIN TRANSACTION
```

```
...  
SELECT ORDER_TOTAL  
FROM ORDERS WHERE CUSTOMER_ID = ?
```

```
...  
SELECT CREDIT_LIMIT  
FROM CUSTOMERS WHERE CUSTOMER_ID = ?
```

```
...  
INSERT INTO ORDERS ...
```

```
...  
COMMIT TRANSACTION
```

# 2PC is not a viable option

- Guarantees consistency

**BUT**

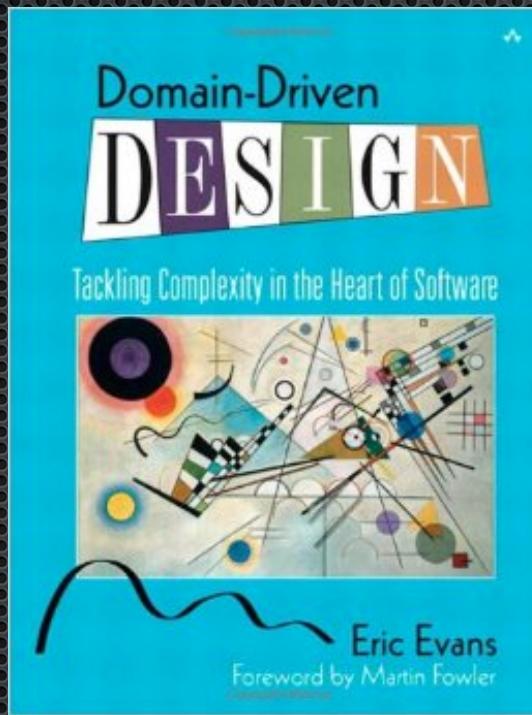
- 2PC is best avoided
- Not supported by many NoSQL databases etc.
- CAP theorem  $\Rightarrow$  2PC impacts availability
- ....

Doesn't fit with the  
NoSQL DB “transaction” model

# Agenda

- The problem with Domain Models and microservices
- Overview of aggregates
- Maintaining consistency between aggregates
- Using event sourcing with Aggregates
- Example application

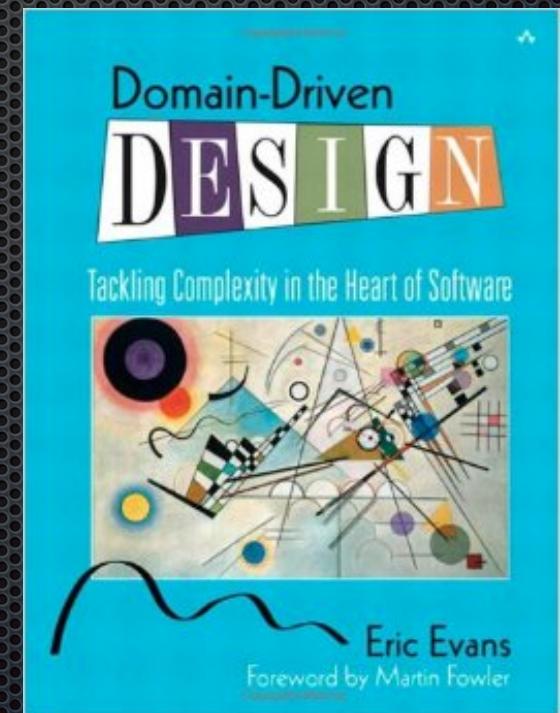
# Domain Driven Design: read it!



# Domain Driven Design - building blocks

- ❖ Entity
- ❖ Value object
- ❖ Services
- ❖ Repositories
- ❖ Aggregates

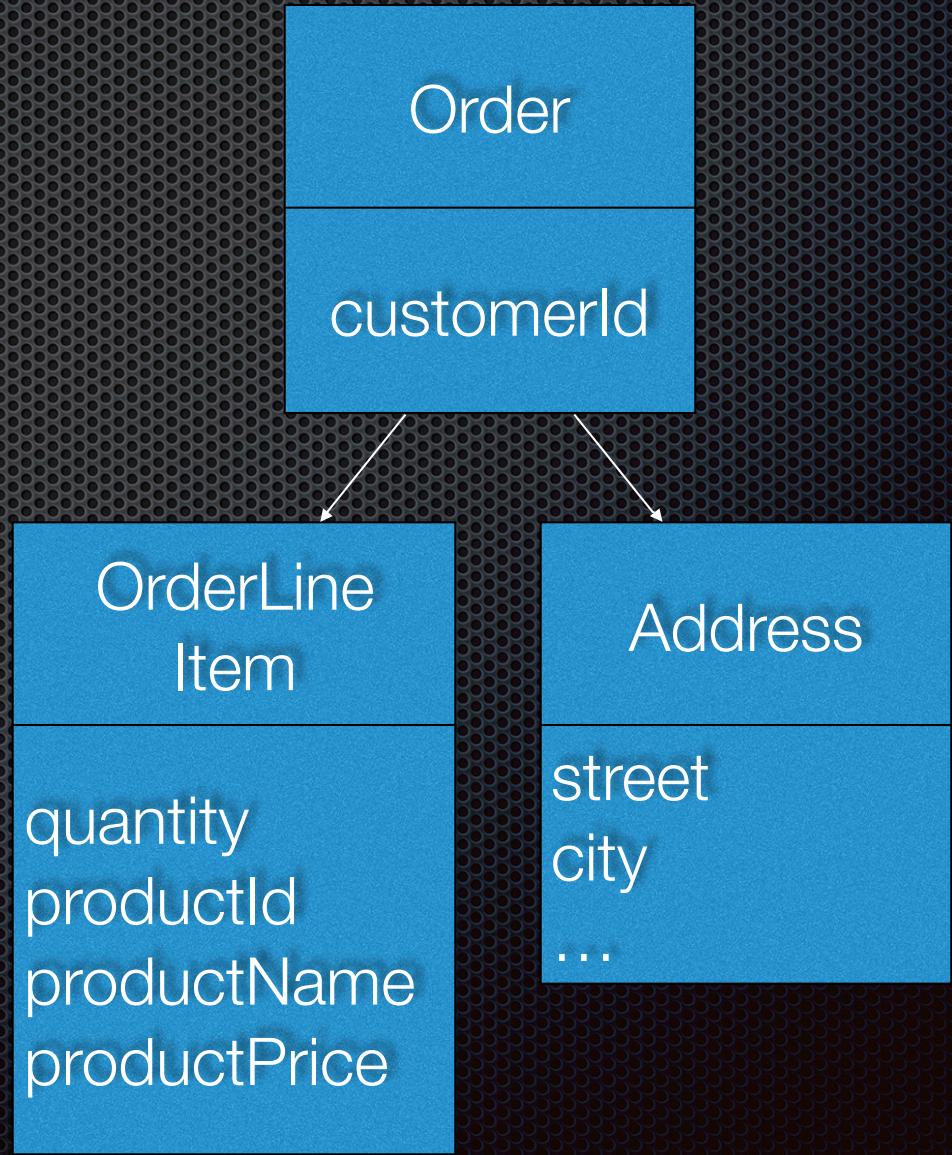
Adopted



Ignored  
(at least by me)

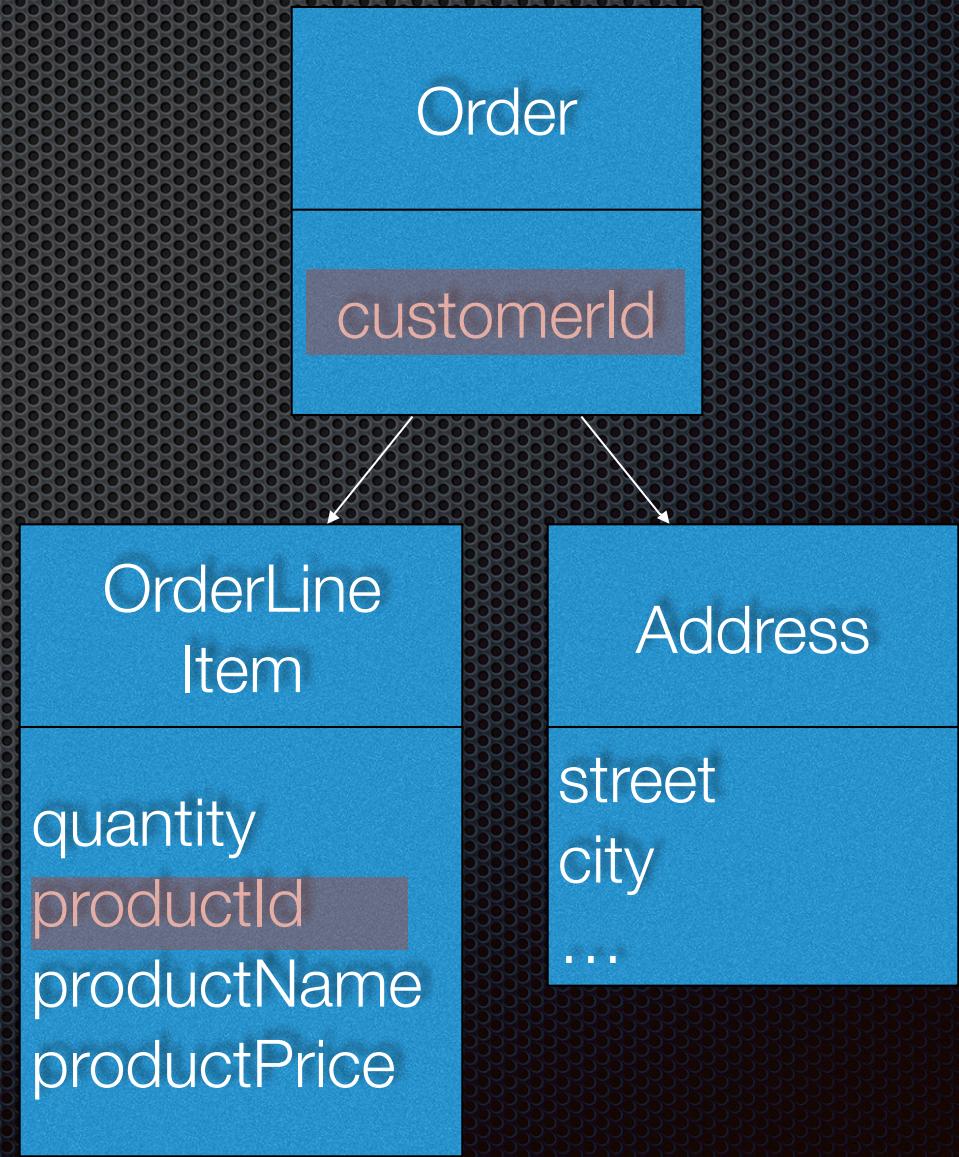
# About Aggregates

- Cluster of objects that can be treated as a unit
- Graph consisting of a root entity and one or more other entities and value objects
- Typically business entities are Aggregates, e.g. customer, Account, Order, Product, ...



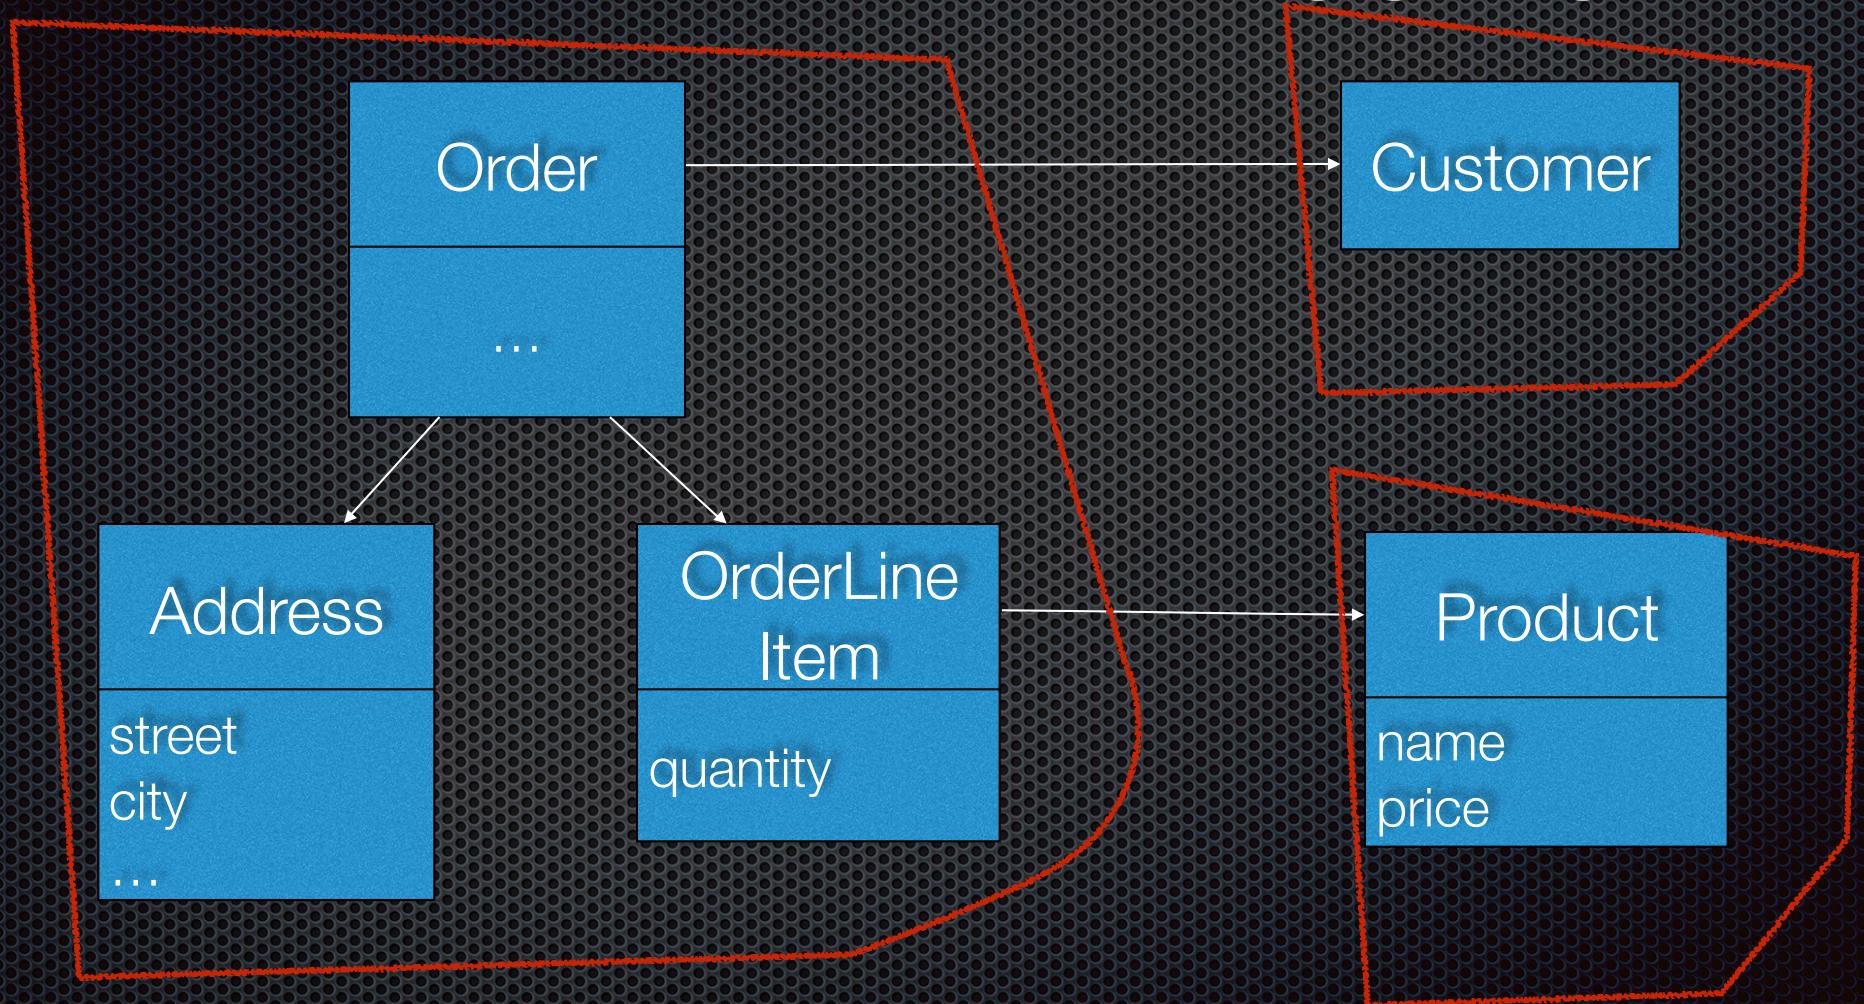
# Aggregate: rule #1

Reference other  
aggregate roots via  
**identity**  
**(primary key)**



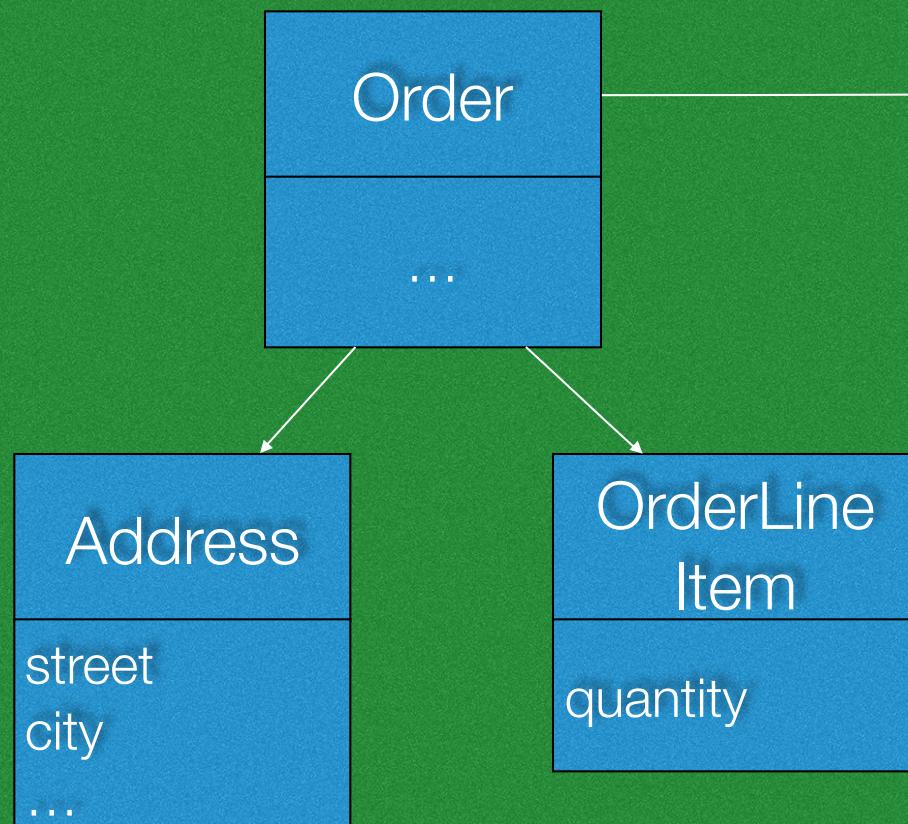
# Foreign keys in a Domain Model?!?

# Domain model = collection of **loosely** connected aggregates

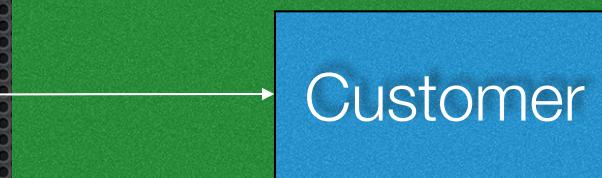


# Easily partition into microservices

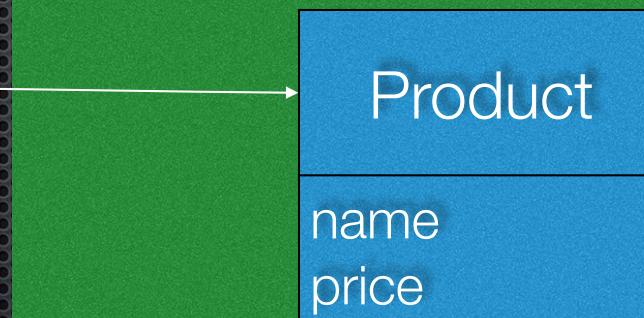
Order service



Customer service



Product service



# Aggregate rule #2

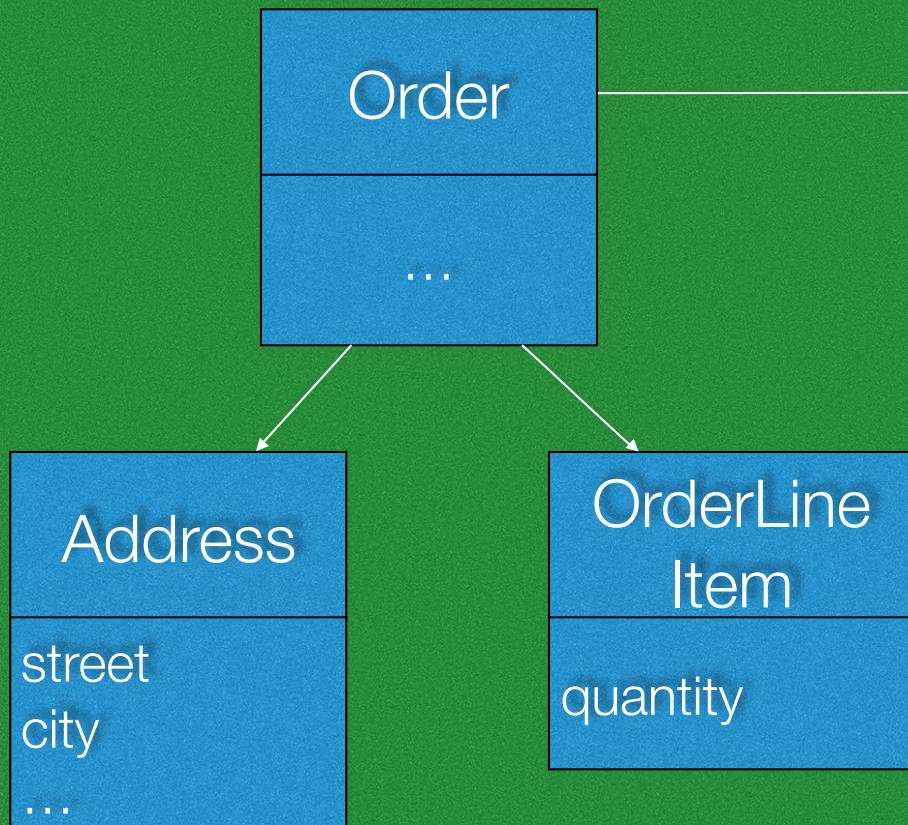
Transaction

=

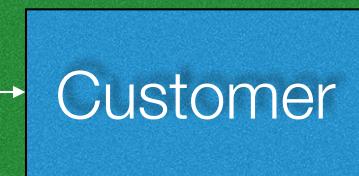
processing **one** command by  
**one** aggregate

# Transaction scope = service

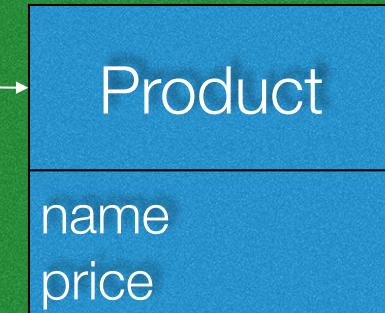
Order service



Customer service



Product service



Transaction scope

=

NoSQL database  
“transaction”

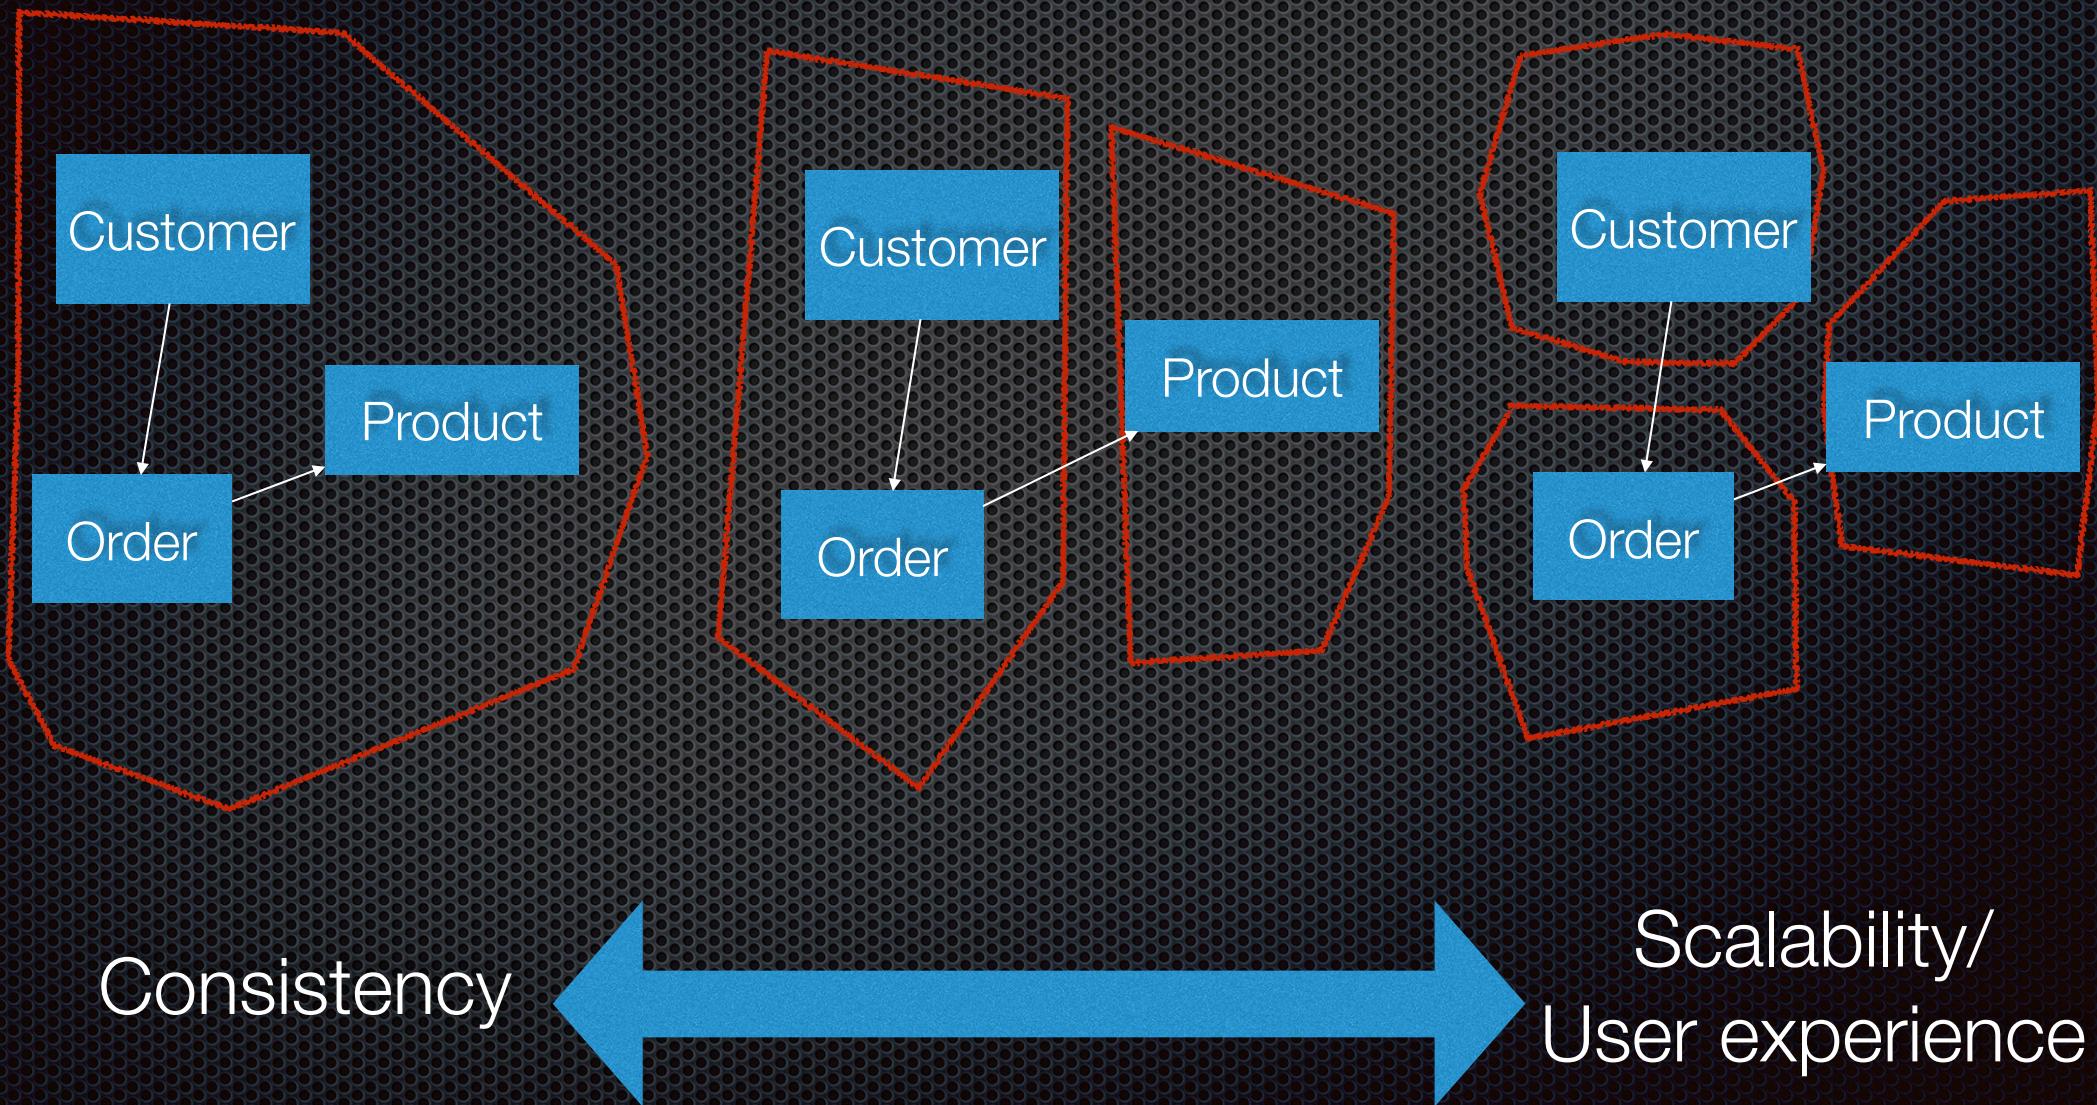
# Aggregate granularity

- If an update must be atomic then it must be handled by a single aggregate

**Therefore**

- Aggregate granularity is important

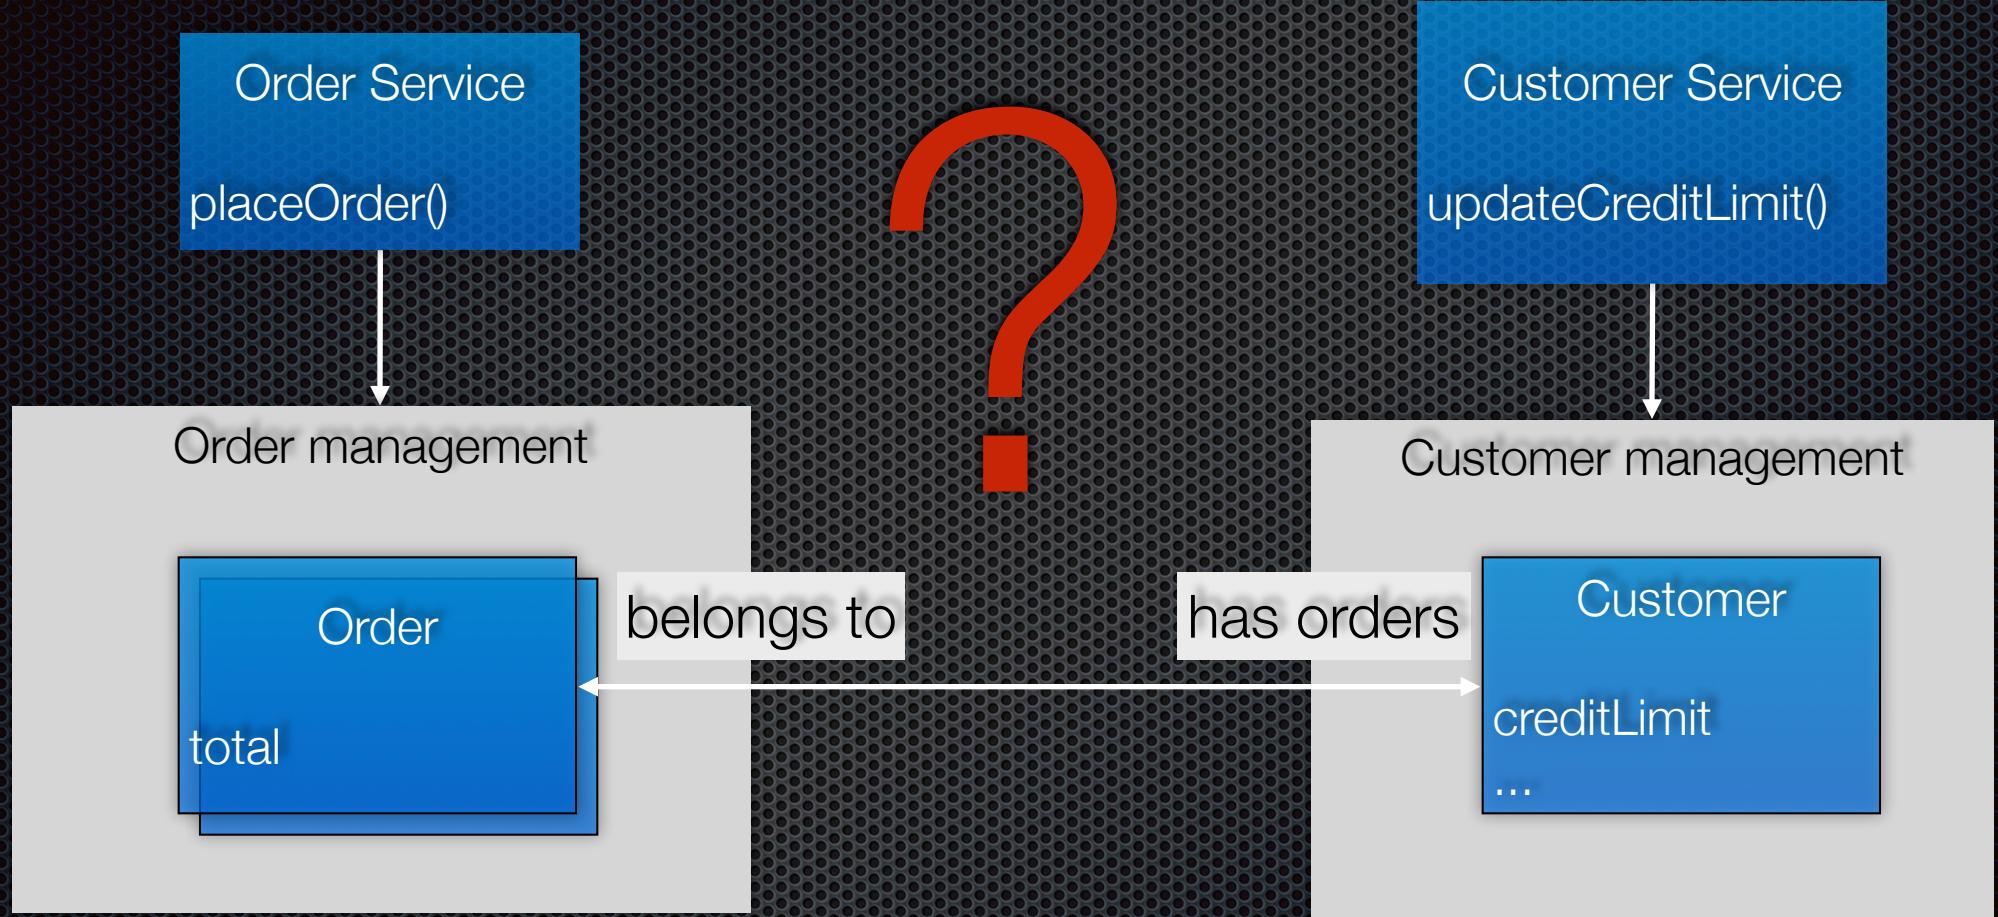
# Aggregate granularity



# Agenda

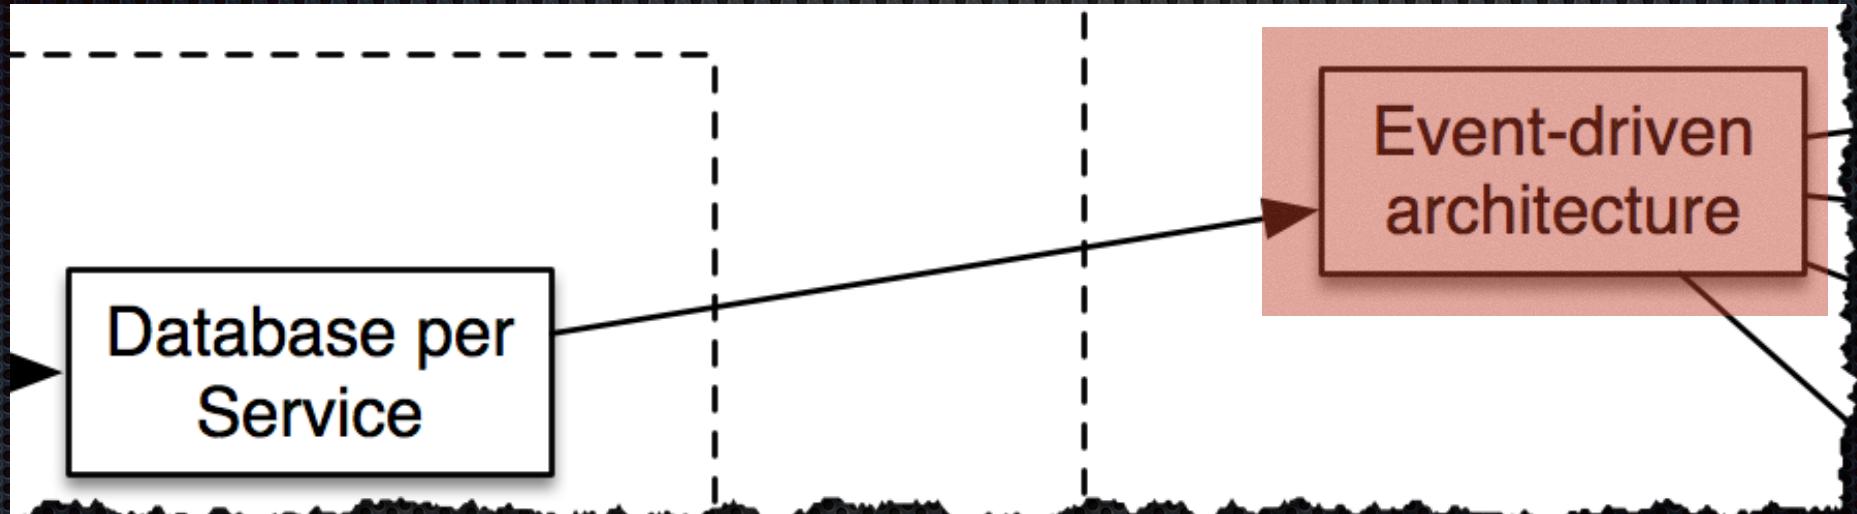
- The problem with Domain Models and microservices
- Overview of aggregates
- Maintaining consistency between aggregates
- Using event sourcing with Aggregates
- Example application

# How to maintain data consistency between aggregates?

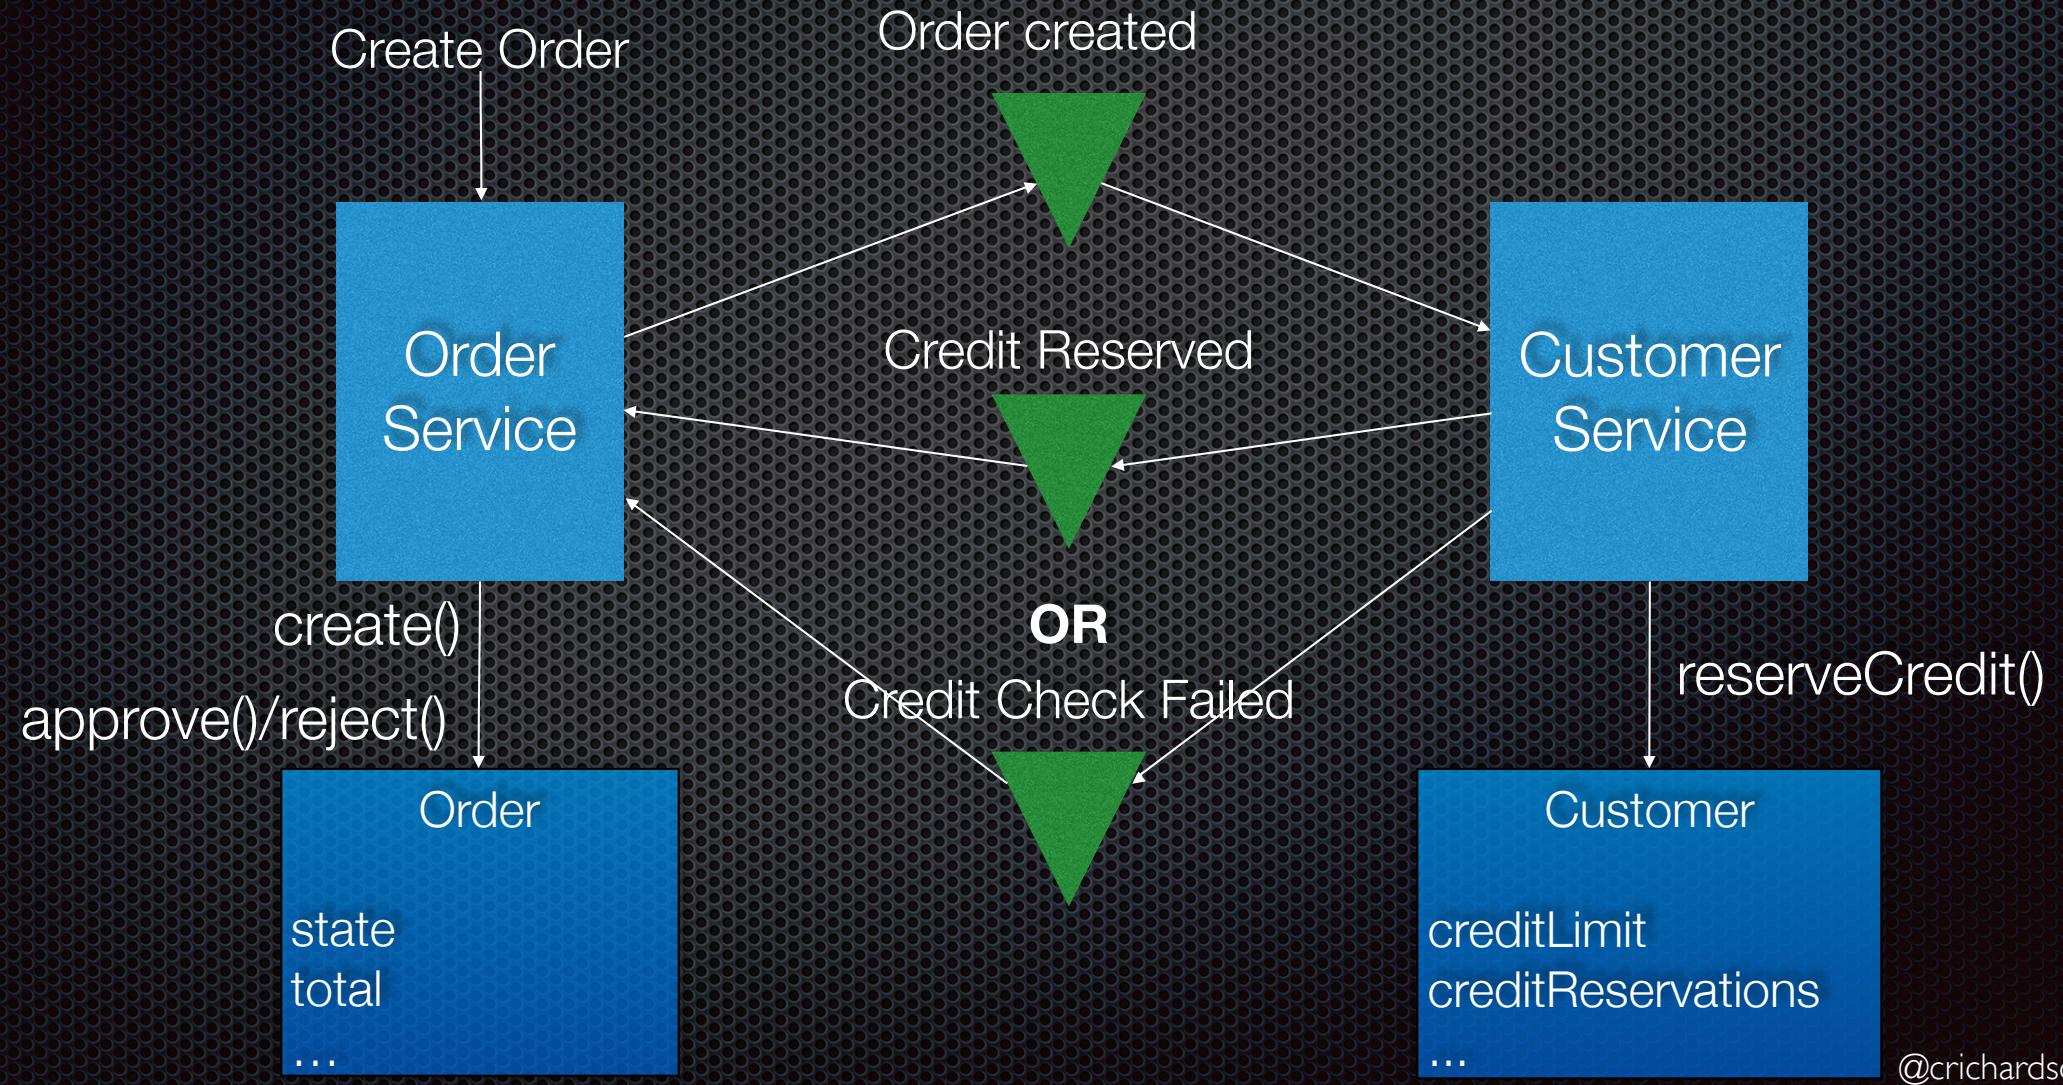


Invariant:  
 $\text{sum(open order.total)} \leq \text{customer.creditLimit}$

# Event-driven architecture



# Use event-driven, eventually consistent order processing



# Eventually consistent credit checking

createOrder()



Order Management

Order  
id : 4567  
total: 343  
state = APPROVED

Customer Management

Customer  
creditLimit : 12000  
creditReservations: { 4567 -> 343}

**Subscribes to:**  
CreditReservedEvent

**publishes:**

OrderCreatedEvent

**Subscribes to:**  
OrderCreatedEvent

**Publishes:**

CreditReservedEvent

Message Bus

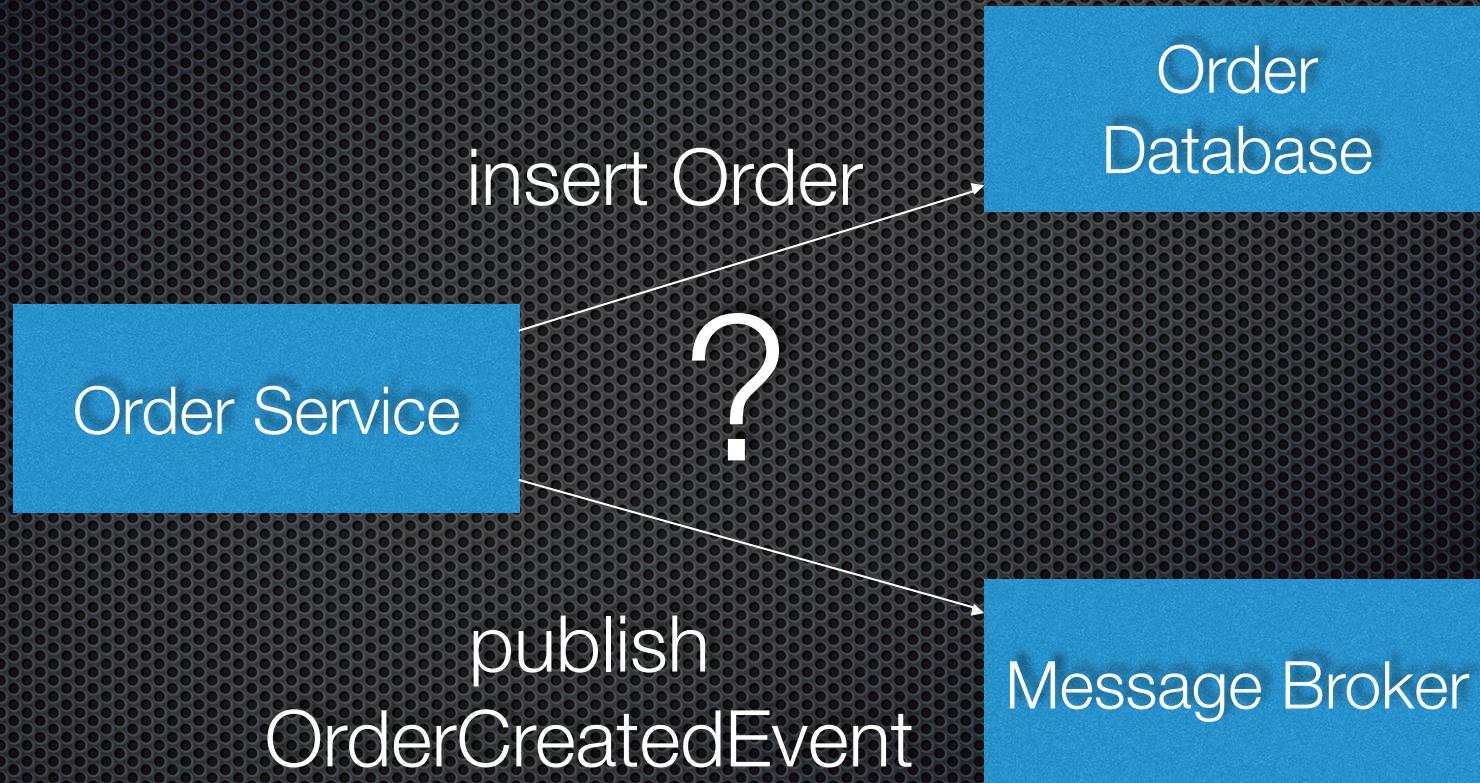
# Complexity of compensating transactions

- ACID transactions can simply rollback

**BUT**

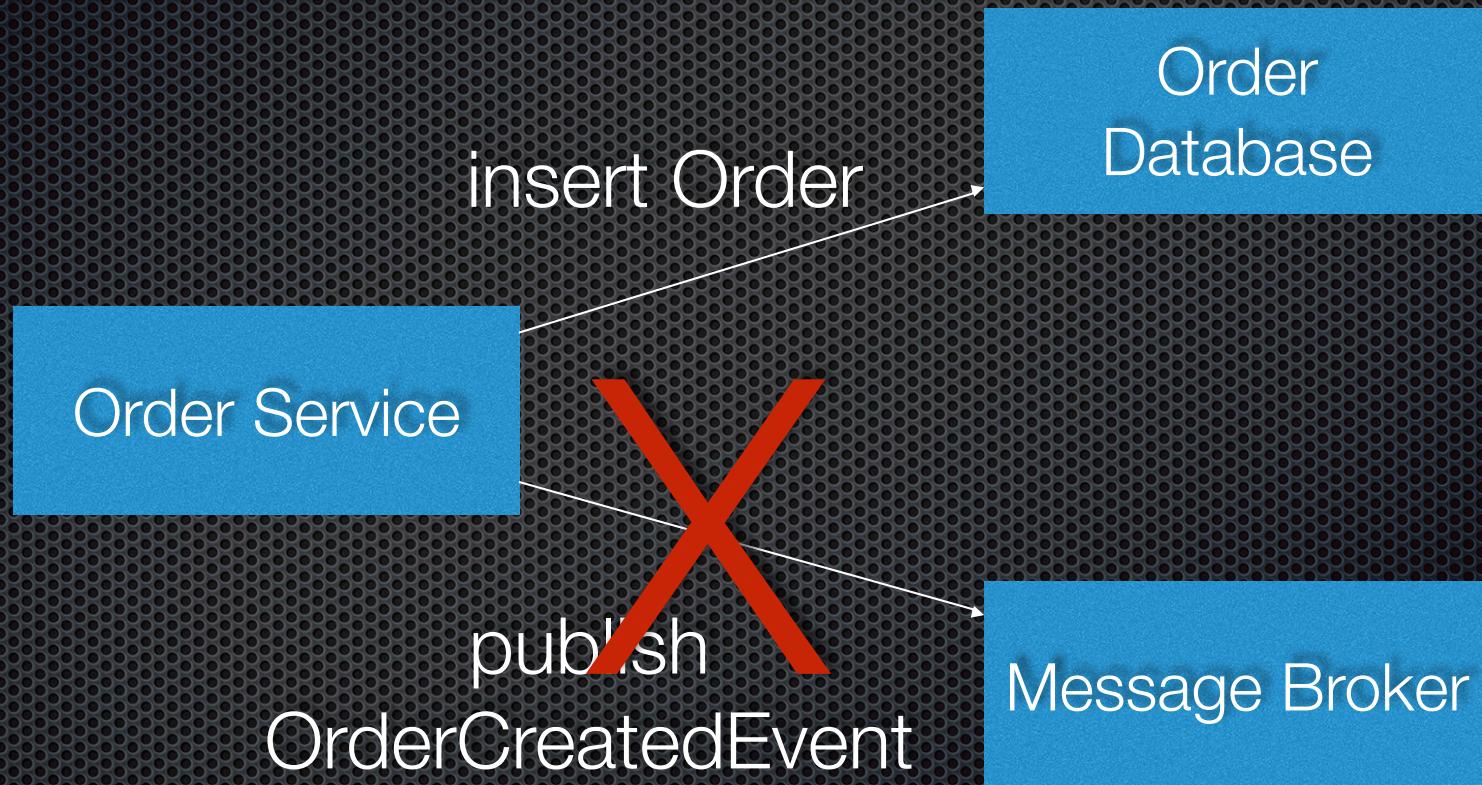
- Developer must write application logic to “rollback” eventually consistent transactions
- For example:
  - CreditCheckFailed => cancel Order
  - Money transfer destination account closed => credit source account
- Careful design required!

# How atomically update database and publish an event



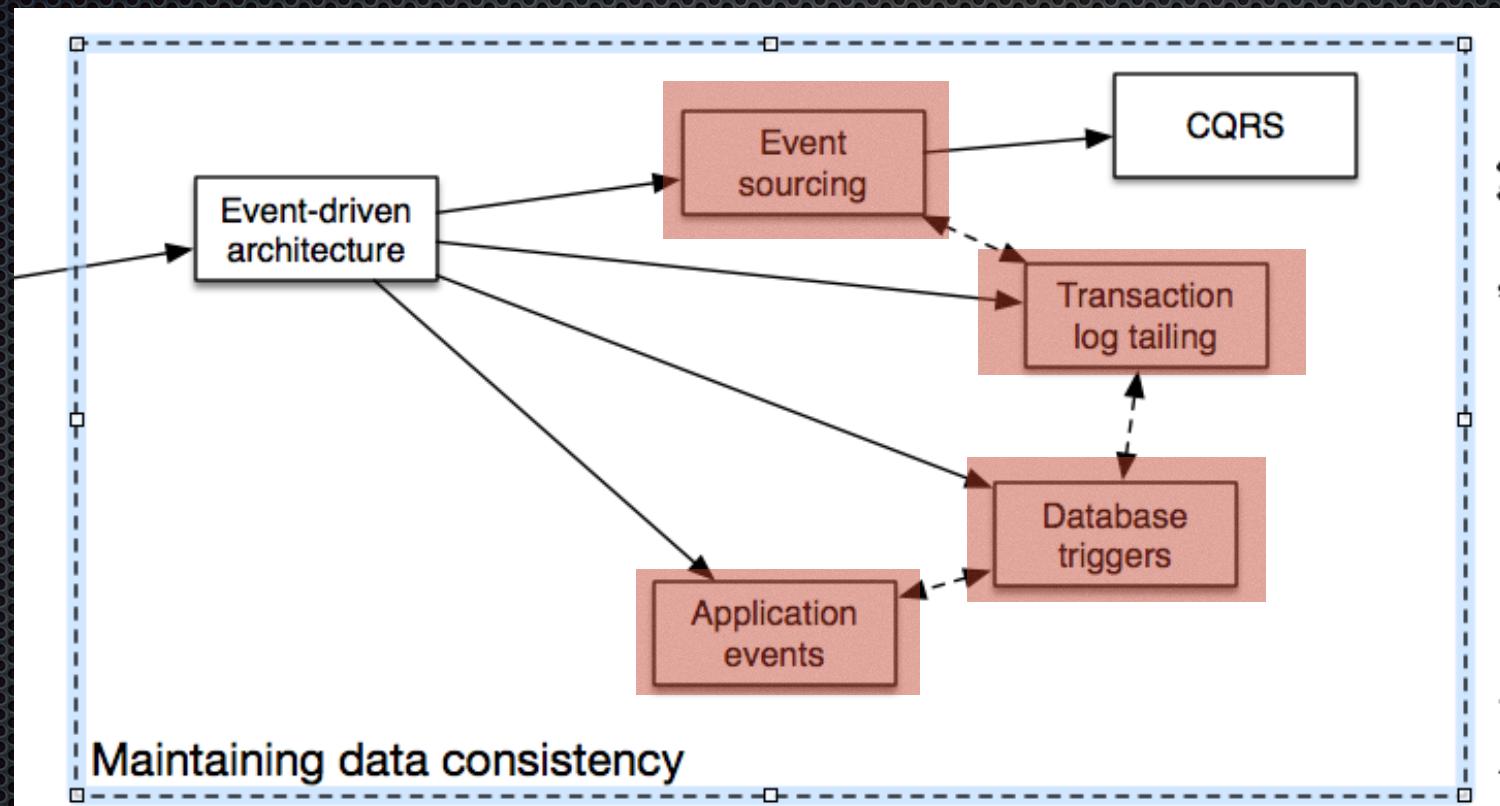
dual write problem

# Failure = inconsistent system



# 2PC is not an option

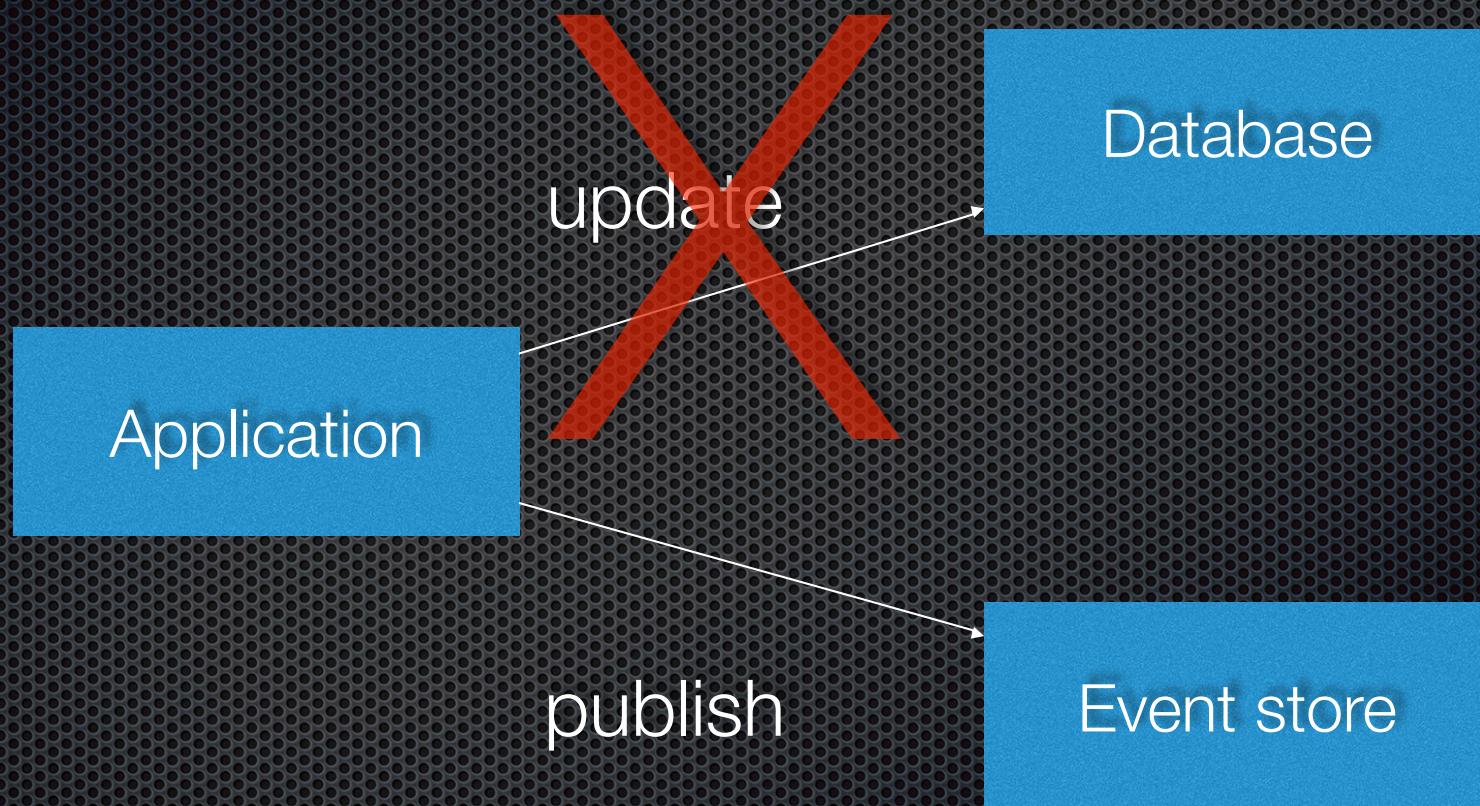
# Reliably publish events when state changes



# Agenda

- The problem with Domain Models and microservices
- Overview of aggregates
- Maintaining consistency between aggregates
- Using event sourcing with Aggregates
- Example application

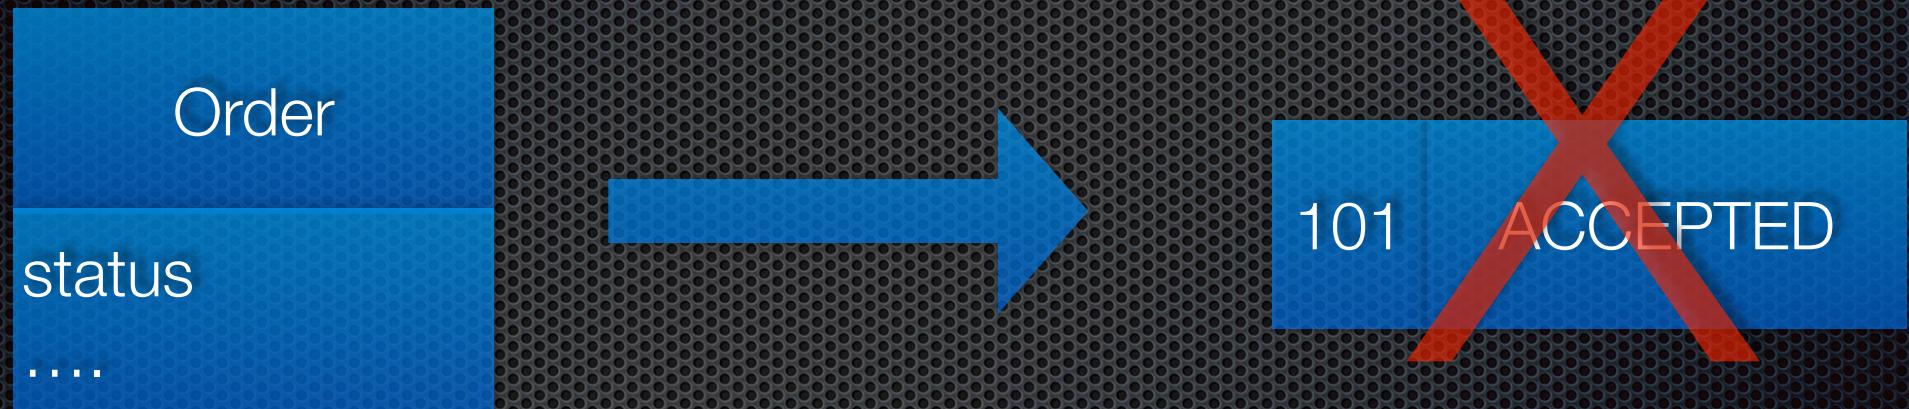
# Event sourcing = event-centric persistence



# Event sourcing

- For each domain object (i.e. DDD aggregate):
  - Identify (state changing) domain events, e.g. use Event Storming
  - Define Event classes
- For example, Order events: OrderCreated, OrderCancelled, OrderApproved, OrderRejected, OrderShipped

# Persists events NOT current state



# Persists events NOT current state

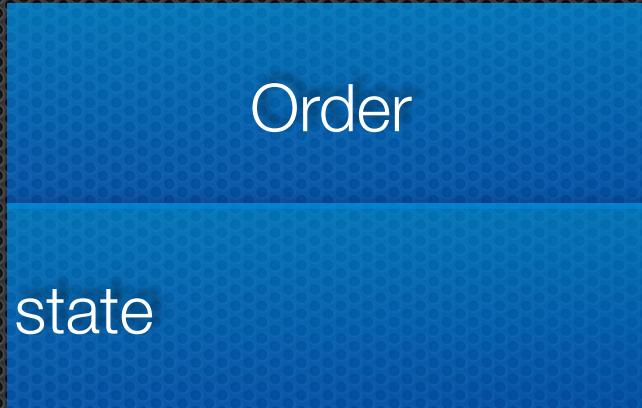
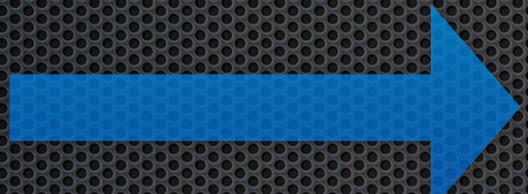
Event table

Entity id	Entity type	Event id	Event type	Event data
101	Order	901	OrderCreated	...
101	Order	902	OrderApproved	...
101	Order	903	OrderShipped	...

# Replay events to recreate state

Events

OrderCreated(...)  
OrderAccepted(...)  
OrderShipped(...)

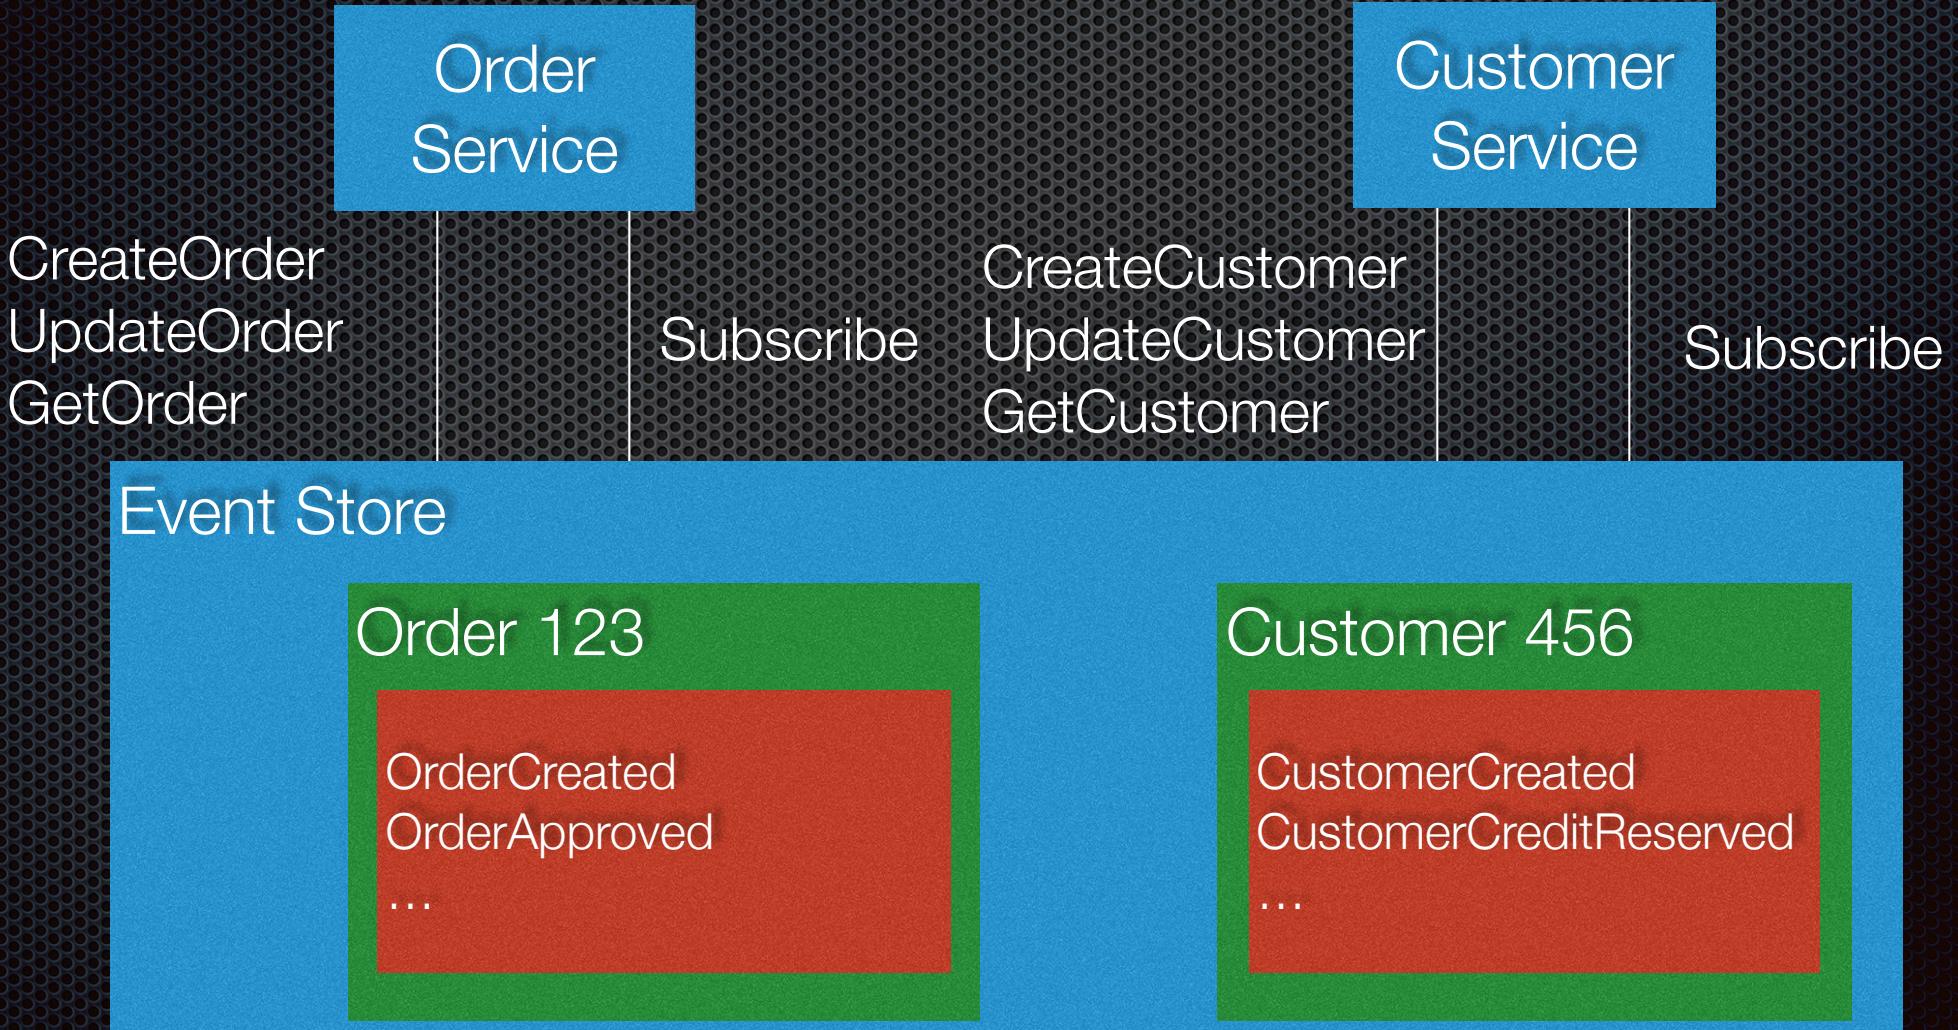


Periodically snapshot to avoid loading all events

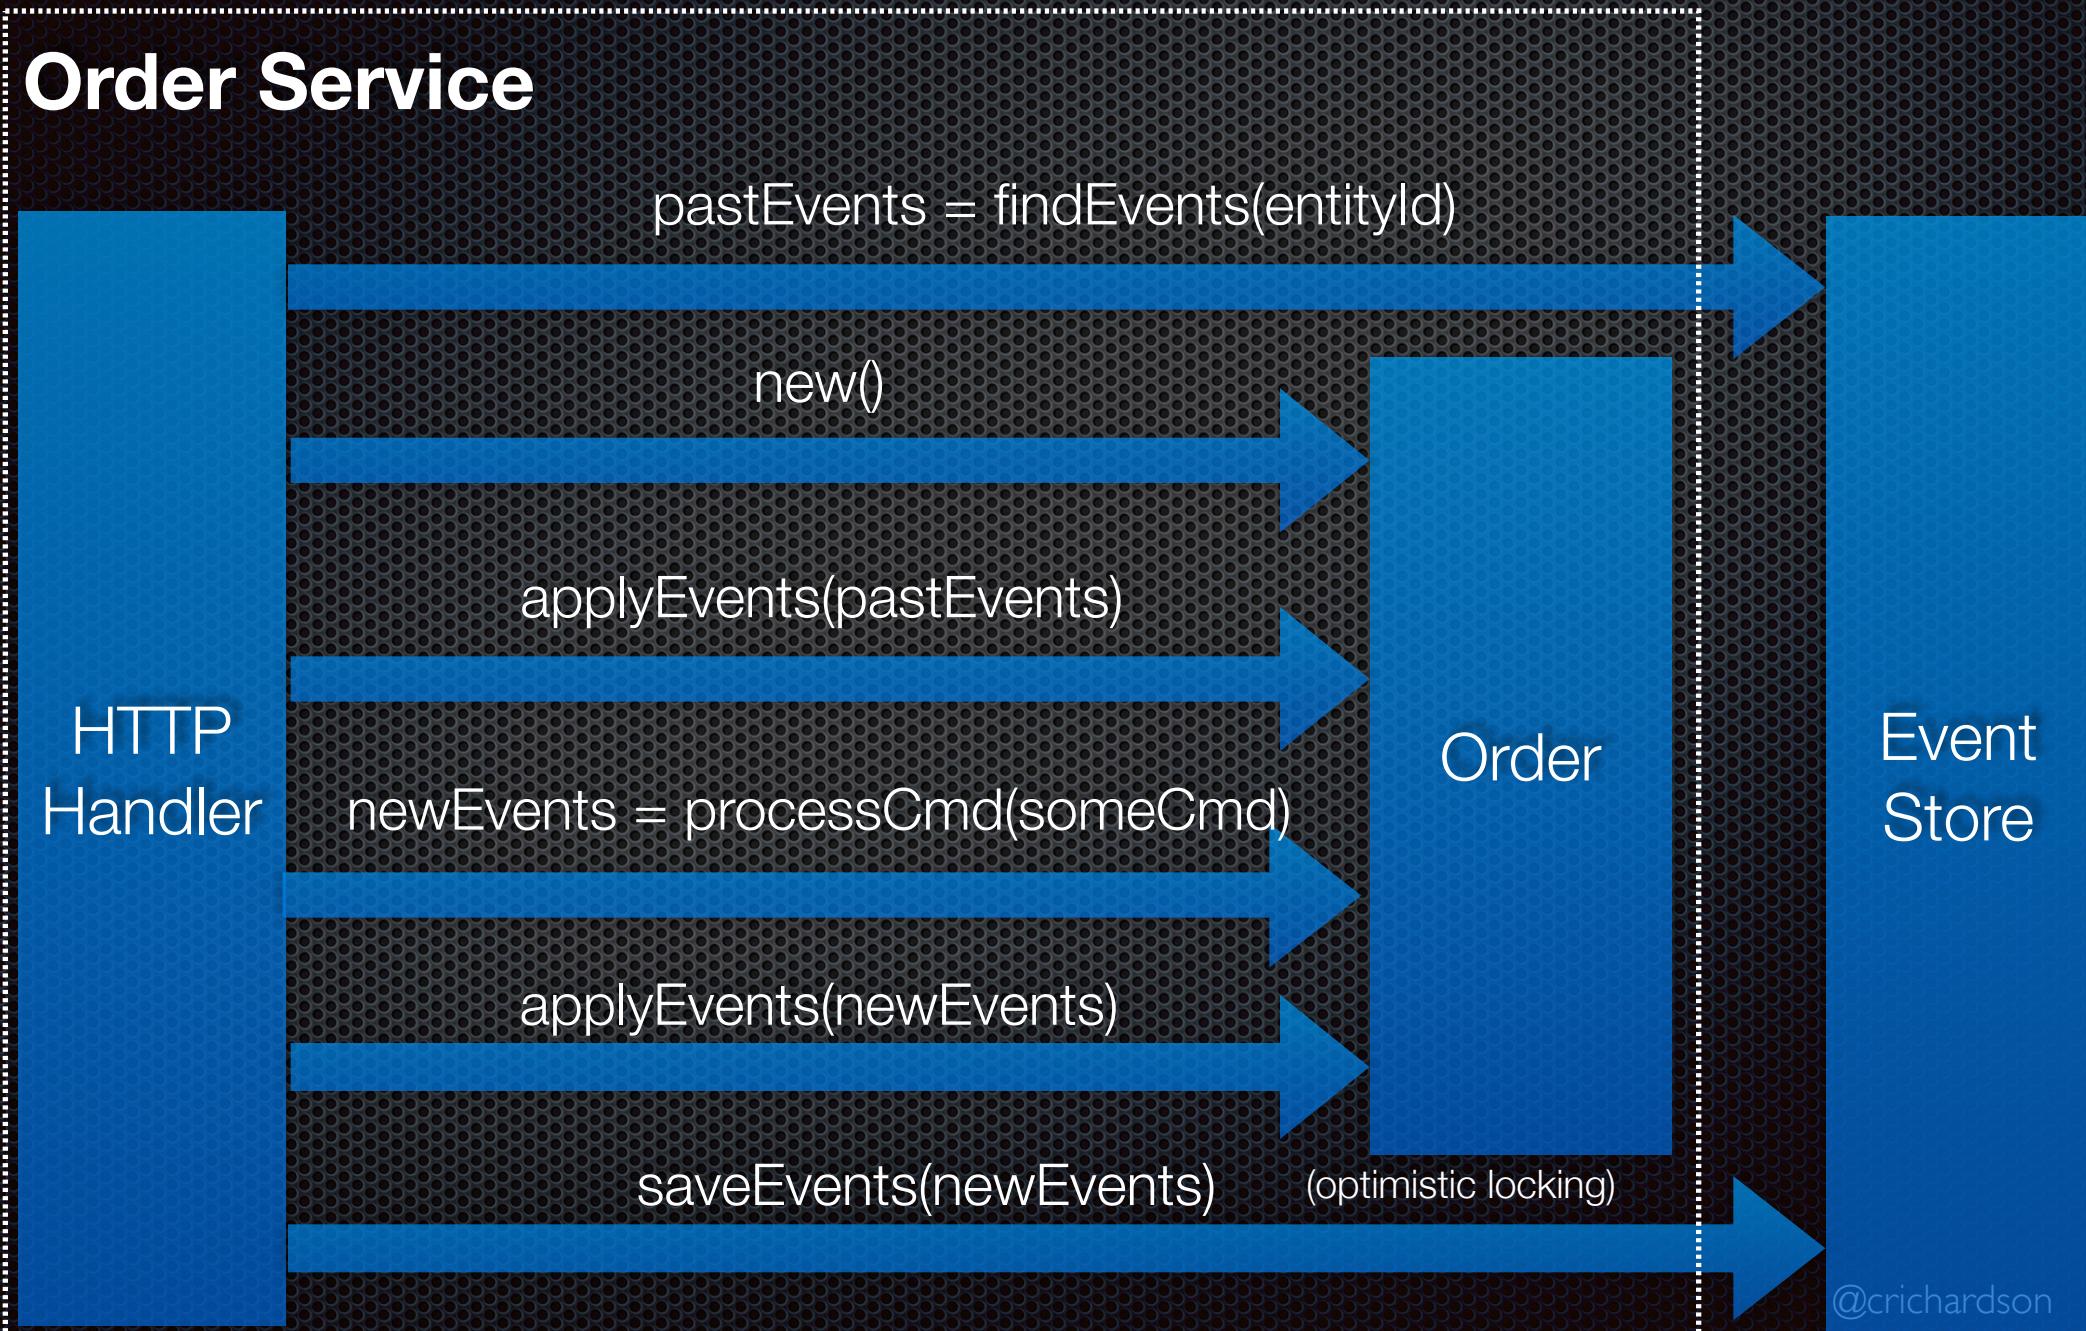
The present is a fold over  
history

currentState = foldl(applyEvent, initialState, events)

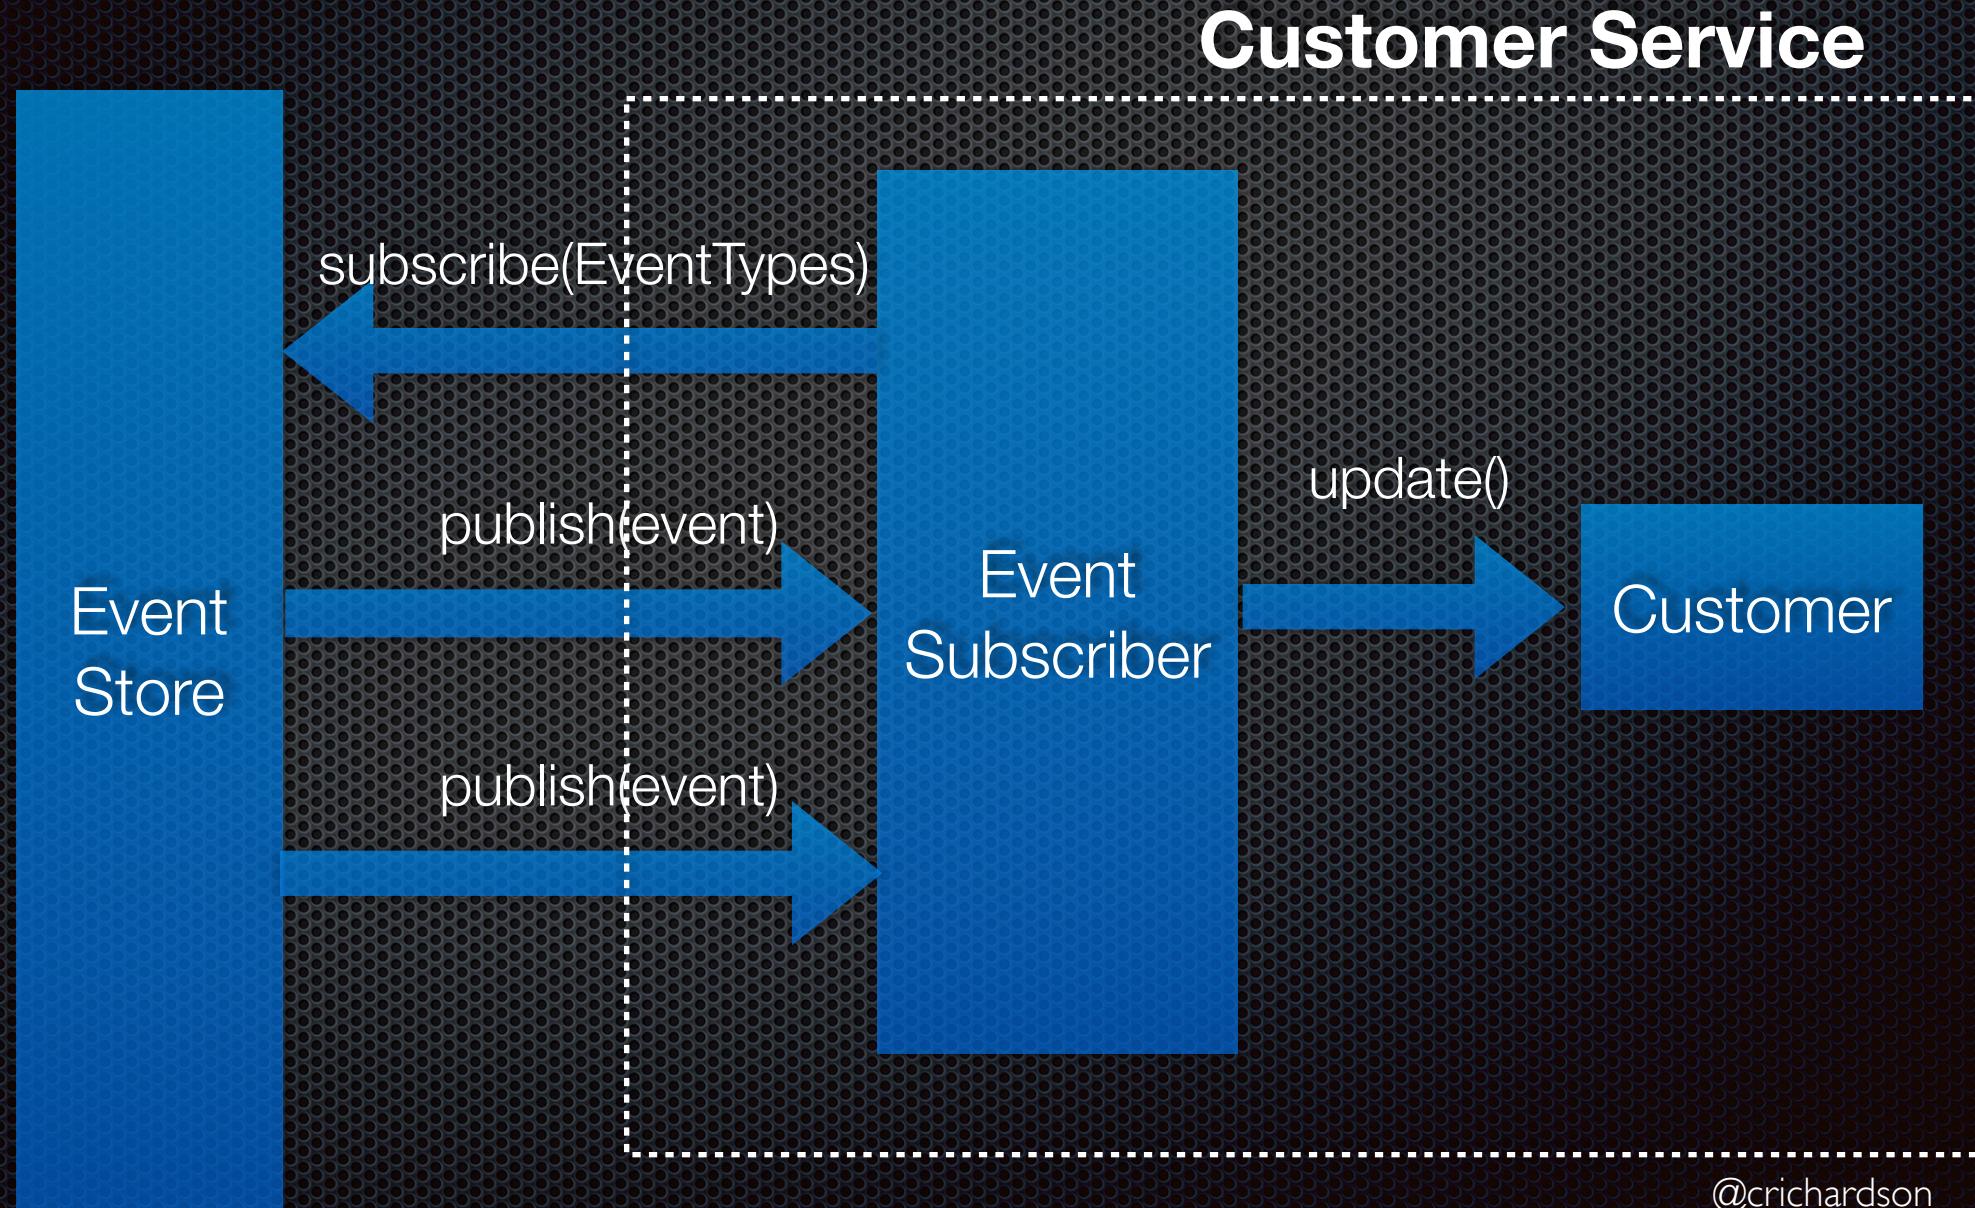
# Application architecture



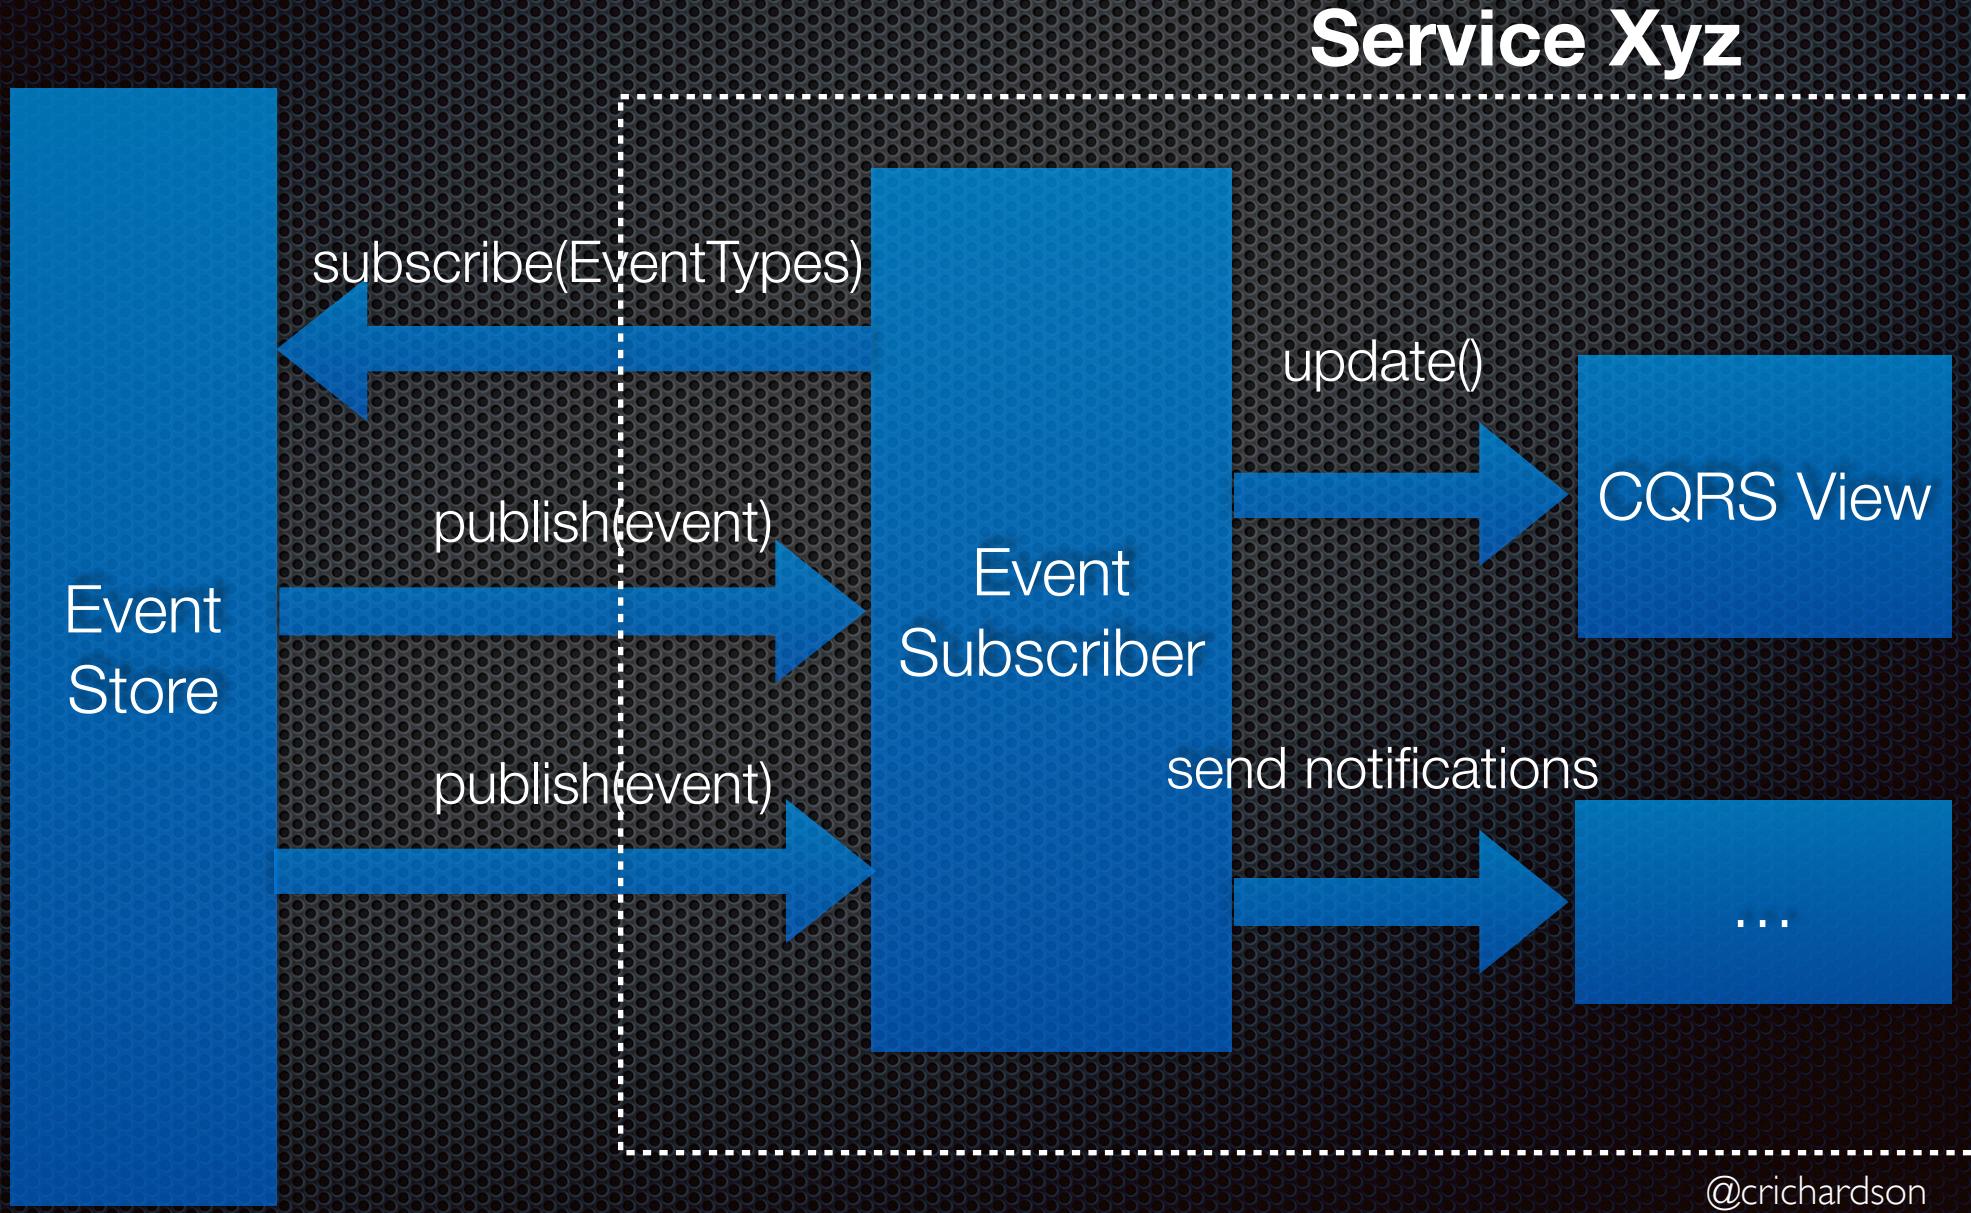
# Request handling in an event sourced application



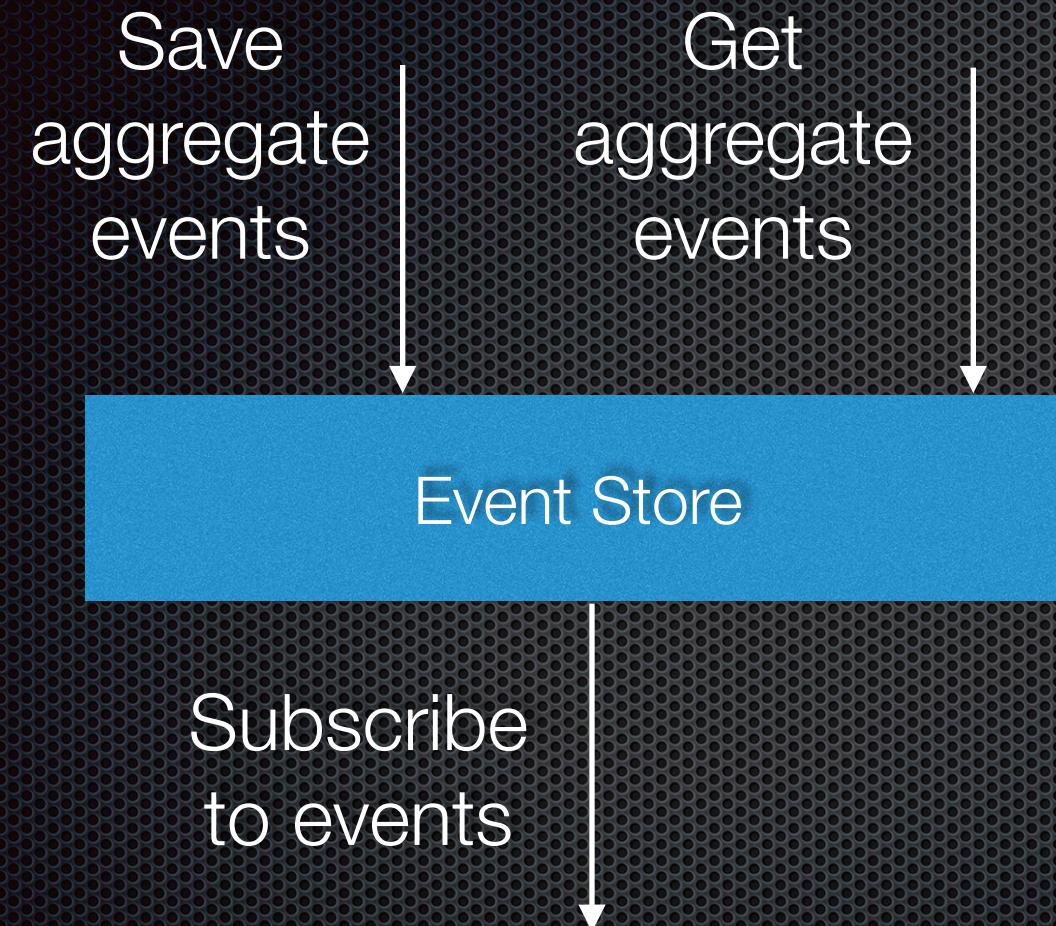
# Event Store publishes events consumed by other services



# Event Store publishes events consumed by other services



# Event store = database + message broker



- Hybrid database and message broker
- Implementations:
  - Home grown/DIY
  - [geteventstore.com](http://geteventstore.com) by Greg Young
  - <http://eventuate.io> (mine)

# Benefits of event sourcing

- Solves data consistency issues in a Microservice/NoSQL based architecture
- Reliable event publishing: publishes events needed by predictive analytics etc, user notifications,...
- Eliminates O/R mapping problem (mostly)
- Reifies state changes:
  - Built in, reliable audit log
  - temporal queries
- Preserved history ⇒ More easily implement future requirements

# Drawbacks of event sourcing...

- Requires application rewrite
- Weird and unfamiliar style of programming
- Events = a historical record of your bad design decisions
- Must detect and ignore duplicate events
  - Idempotent event handlers
  - Track most recent event and ignore older ones
  - ...

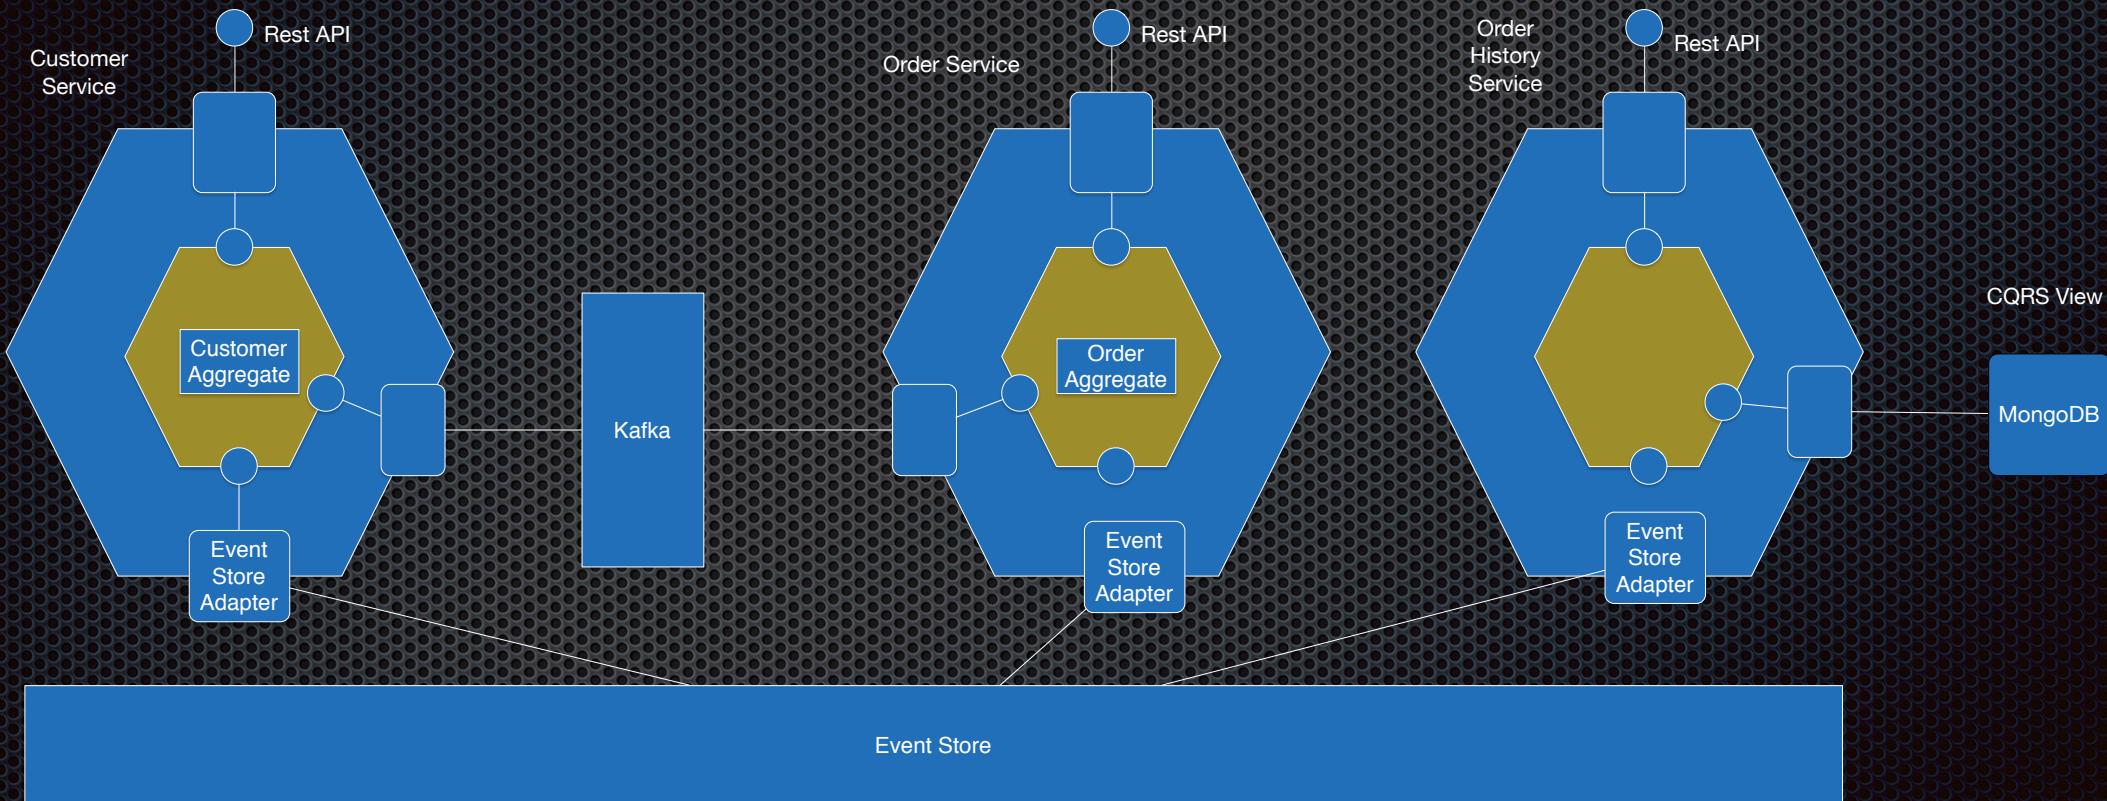
# .... Drawbacks of event sourcing

- Querying the event store can be challenging
- Some queries might be complex/inefficient, e.g. accounts with a balance > X
- Event store might only support lookup of events by entity id
- Must use Command Query Responsibility Segregation (CQRS) to handle queries ⇒ application must handle eventually consistent data

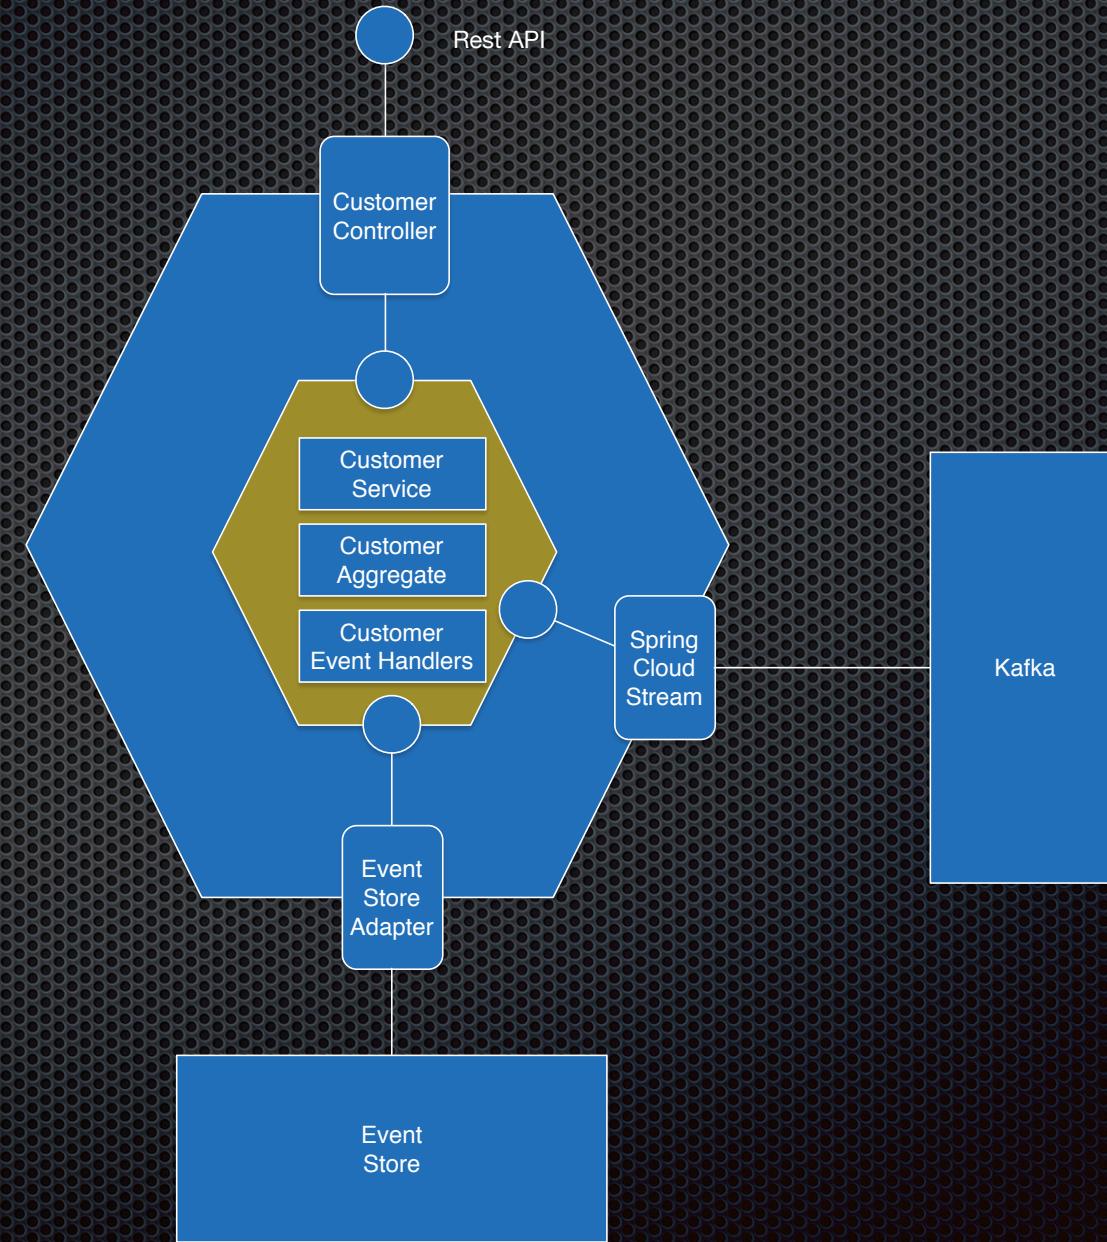
# Agenda

- The problem with Domain Models and microservices
- Overview of aggregates
- Maintaining consistency between aggregates
- Using event sourcing with Aggregates
- Example application

# Orders and Customers example



# Customer microservice



# The Customer aggregate

State

Customer

Money creditLimit

Map<OrderId, Money> creditReservations

List<Event> process(CreateCustomerCommand cmd) { ... }

List<Event> process(ReserveCreditCommand cmd) { ... }

...

void apply(CustomerCreatedEvent anEvent) { ... }

void apply(CreditReservedEvent anEvent) { ... }

...

Behavior

# Familiar concepts restructured

```
class Customer {  
  
    public void reserveCredit(  
        orderId : String,  
        amount : Money) {  
  
        // verify  
  
        // update state  
        this.xyz = ...  
    }  
}
```



```
public List<Event> process(  
    ReserveCreditCommand cmd) {  
  
    // verify  
    ...  
    return singletonList(  
        new CreditReservedEvent(...);  
    )  
}
```



```
public void apply(  
    CreditReservedEvent event) {  
  
    // update state  
    this.xyz = event.xyz  
}
```

# Customer command processing

```
public class Customer extends ReflectiveMutableCommandProcessingAggregate<Customer, CustomerCommand> {

    private Money creditLimit;
    private Map<EntityIdentifier, Money> creditReservations;

    Money availableCredit() {
        return creditLimit.subtract(creditReservations.values().stream().reduce(Money.ZERO, Money::add));
    }

    Money getCreditLimit() { return creditLimit; }

    public List<Event> process(CreateCustomerCommand cmd) {
        return EventUtil.events(new CustomerCreatedEvent(cmd.getName(), cmd.getCreditLimit()));
    }

    public List<Event> process(ReserveCreditCommand cmd) {
        if (availableCredit().isGreaterThanOrEqualTo(cmd.getOrderTotal()))
            return EventUtil.events(new CustomerCreditReservedEvent(cmd.getOrderId(), cmd.getOrderTotal()));
        else
            return EventUtil.events(new CustomerCreditLimitedExceededEvent(cmd.getOrderId()));
    }

    public void apply(CustomerCreatedEvent event) {...}

    public void apply(CustomerCreditReservedEvent event) {...}

    public void apply(CustomerCreditLimitedExceededEvent event) {...}

}
```

# Customer applying events

```
public class Customer extends ReflectiveMutableCommandProcessingAggregate<Customer, CustomerCommand> {

    private Money creditLimit;
    private Map<EntityIdentifier, Money> creditReservations;

    Money availableCredit() {...}

    Money getCreditLimit() { return creditLimit; }

    public List<Event> process(CreateCustomerCommand cmd) {...}

    public List<Event> process(ReserveCreditCommand cmd) {...}

    public void apply(CustomerCreatedEvent event) {
        this.creditLimit = event.getCreditLimit();
        this.creditReservations = new HashMap<>();
    }

    public void apply(CustomerCreditReservedEvent event) {
        this.creditReservations.put(event.getOrderId(), event.getOrderTotal());
    }

    public void apply(CustomerCreditLimitedExceededEvent event) {
        // Do nothing
    }

}
```

# Customer Controller

```
@RestController
public class CustomerController {

    private CustomerService customerService;

    @Autowired
    public CustomerController(CustomerService customerService) {
        this.customerService = customerService;
    }

    @RequestMapping(value = "/customers", method = RequestMethod.POST)
    public CompletableFuture<CreateCustomerResponse>
            createCustomer(@RequestBody CreateCustomerRequest createCustomerRequest) {
        return customerService.createCustomer(createCustomerRequest.getName(),
                                              createCustomerRequest.getCreditLimit())
            .thenApply(ewidv -> new CreateCustomerResponse(ewidv.getEntityId()));
    }
}
```

# Creating a Customer

```
public class CustomerServiceImpl implements CustomerService {  
  
    private final AggregateRepository<Customer, CustomerCommand> customerRepository;  
  
    public CustomerServiceImpl(AggregateRepository<Customer, CustomerCommand> customerRepository) {  
        this.customerRepository = customerRepository;  
    }  
  
    @Override  
    public CompletableFuture<EntityWithIdAndVersion<Customer>>  
        createCustomer(String name, Money creditLimit) {  
        return customerRepository.save(new CreateCustomerCommand(name, creditLimit));  
    }  
}
```

save() concisely specifies:

1. Creates Customer aggregate
2. Processes command
3. Applies events
4. Persists events

# Event handling in Customers

Triggers BeanPostProcessor

Durable subscription name

```
@EventSubscriber(id = "customerWorkflow")
public class CustomerWorkflow {

    private Logger logger = LoggerFactory.getLogger(getClass());

    private Source source;

    public CustomerWorkflow(Source source) { this.source = source; }

    @EventHandlerMethod
    public CompletableFuture<?> reserveCredit(EventHandlerContext<OrderCreatedEvent> ctx) {
        OrderCreatedEvent event = ctx.getEvent();
        Money orderTotal = event.getOrderTotal();
        String customerId = event.getCustomerId();
        String orderId = ctx.getEntityId();

        return ctx.update(Customer.class, customerId, new ReserveCreditCommand(orderTotal, orderId))
            .handleAsync()
                .recoverWith(singletonMap(EntityNotFoundException.class, (t) -> {
                    rejectOrder(orderId);
                    return null;
                }));
    }
}
```

1. Load Customer aggregate  
2. Processes command  
3. Applies events  
4. Persists events

Publish message to Spring Cloud Stream (Kafka) when Customer not found

# Summary

- Aggregates are the building blocks of microservices
- Use events to maintain consistency between aggregates
- Event sourcing is a good way to implement a event-driven architecture

• @crichton chris@chrisrichardson.net

A close-up photograph of an emu's head, showing its large, dark brown eyes and hooked beak. The feathers around the eyes are dark and textured. The background is blurred green foliage.

Questions?

<http://learnmicroservices.io>